


Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:
 - transition = (s, a, r, s')
 - difference = $\left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$
 - $Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$
 - Exact Q's
 - Approximate Q's
 - $w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$
- Intuitive interpretation:
 - Adjust weights of active features
 - E.g., if something unexpectedly bad happens, blame the features that were on: disprefer all states with that state's features
- Formal justification: online least squares



Q-learning with Linear Approximation

Algorithm 4: Q-learning with linear approximation.

```
1 Initialize q-value function  $Q$  with random weights  $w$ :  $Q(s, a; w) = \sum_m w_m f_m(s, a)$ ;  
2 for  $episode = 1 \rightarrow M$  do  
3   Get initial state  $s_0$ ;  
4   for  $t = 1 \rightarrow T$  do  
5     With prob.  $\epsilon$ , select a random action  $a_t$ ;  
6     With prob.  $1 - \epsilon$ , select  $a_t \in \operatorname{argmax}_a Q(s_t, a; w)$ ;  
7     Execute selected action  $a_t$  and observe reward  $r_t$  and next state  $s_{t+1}$ ;  
8     Set target  $y_t = \begin{cases} r_t & \text{if episode terminates at step } t + 1; \\ r_t + \gamma \max_{a'} Q(s_{t+1}, a'; w) & \text{otherwise} \end{cases}$ ;  
9     Perform a gradient descent step to update  $w$ :  $w_m \leftarrow w_m + \alpha [y_t - Q(s_t, a_t; w)] f_m(s, a)$ ;
```