
CIS 471/571 (Fall 2023): Introduction to Artificial Intelligence

Lecture 2: Uninformed Search

Thanh H. Nguyen

Most slides are by Pieter Abbeel, Dan Klein, Luke Zettlemoyer, John DeNero,
Stuart Russell, Andrew Moore, or Daniel Lowd
Source: <http://ai.berkeley.edu/home.html>

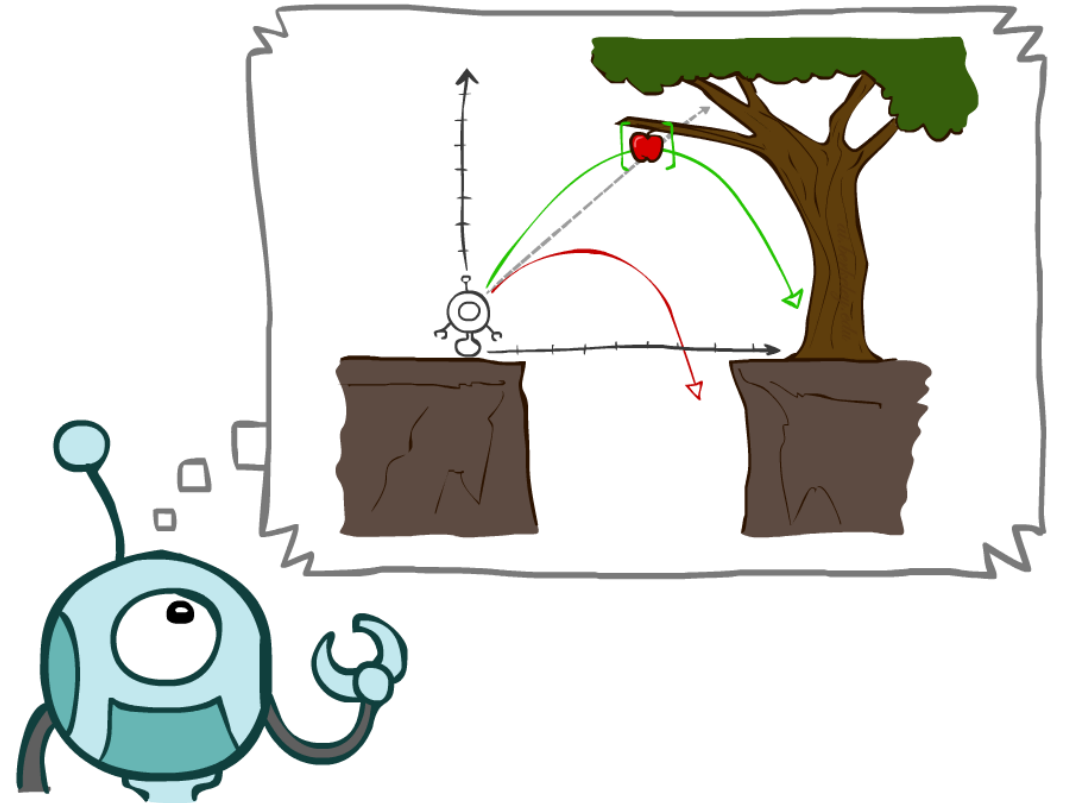


Announcement

- Project 1
 - Deadline: Oct 16th, 2023
- Written Assignment 1
 - Will be posted tomorrow
 - Deadline: Oct 11th, 2023

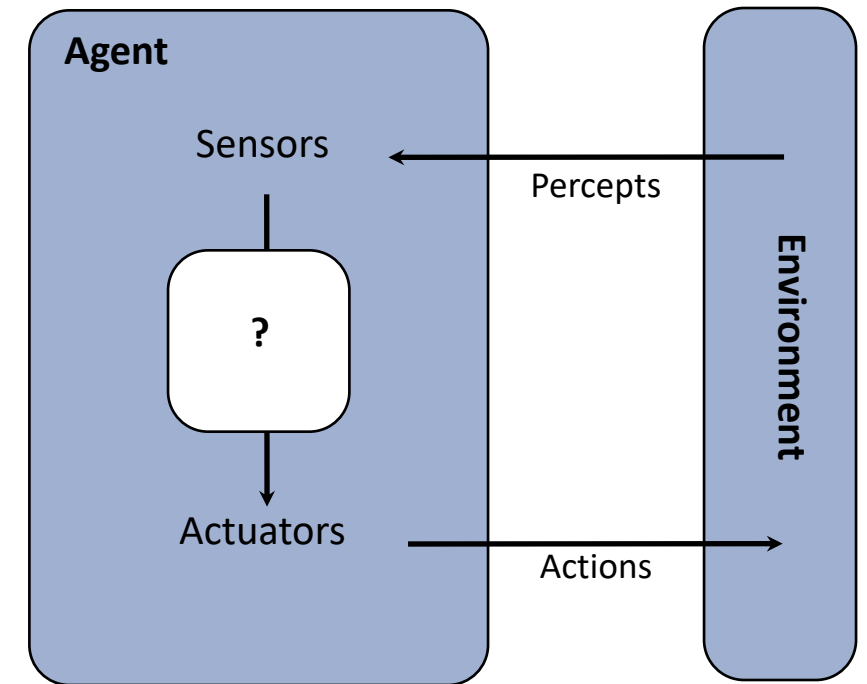
Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
 - Depth-First Search
 - Breadth-First Search
 - Uniform-Cost Search



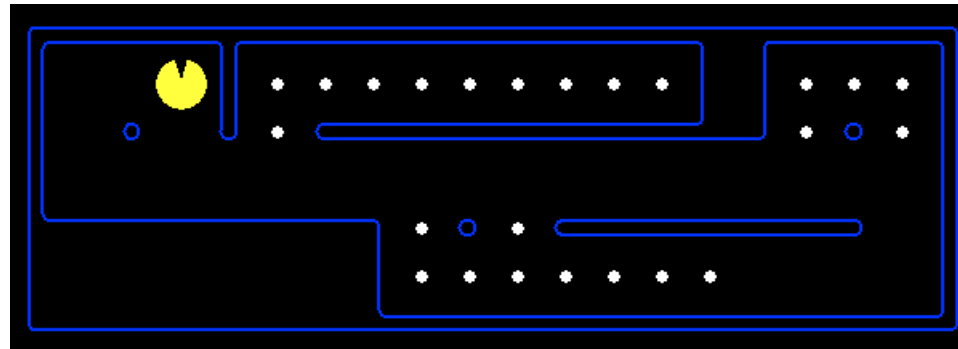
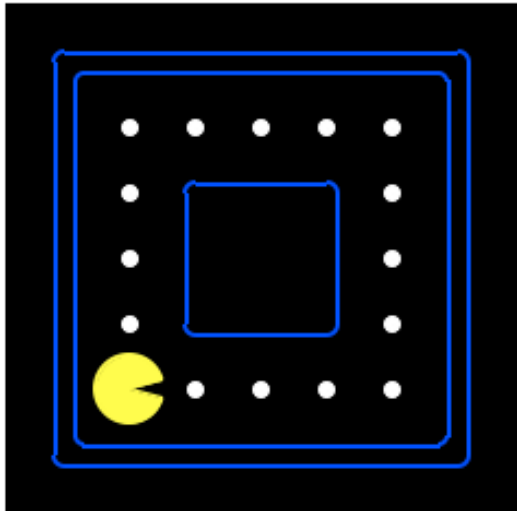
Rational Agents

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its **utility function**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions.

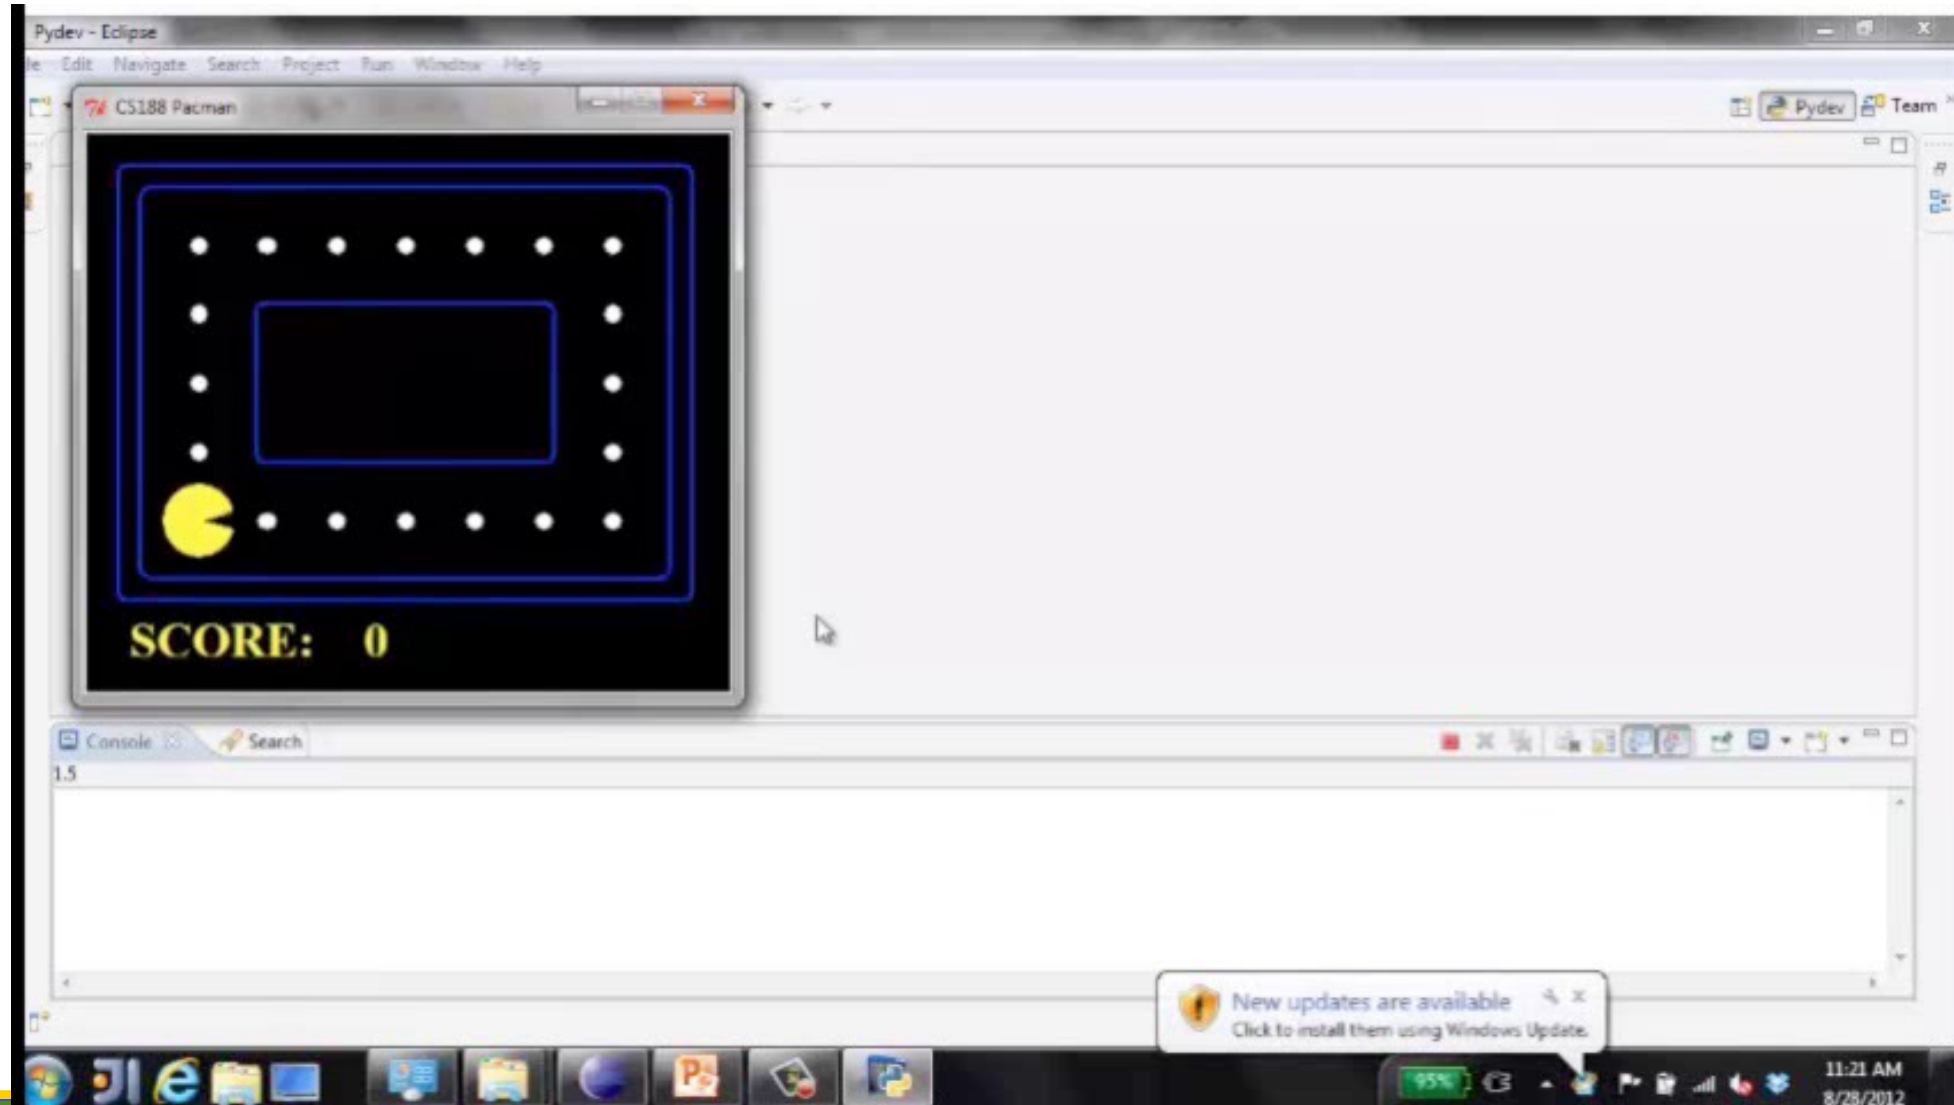


Reflex Agents

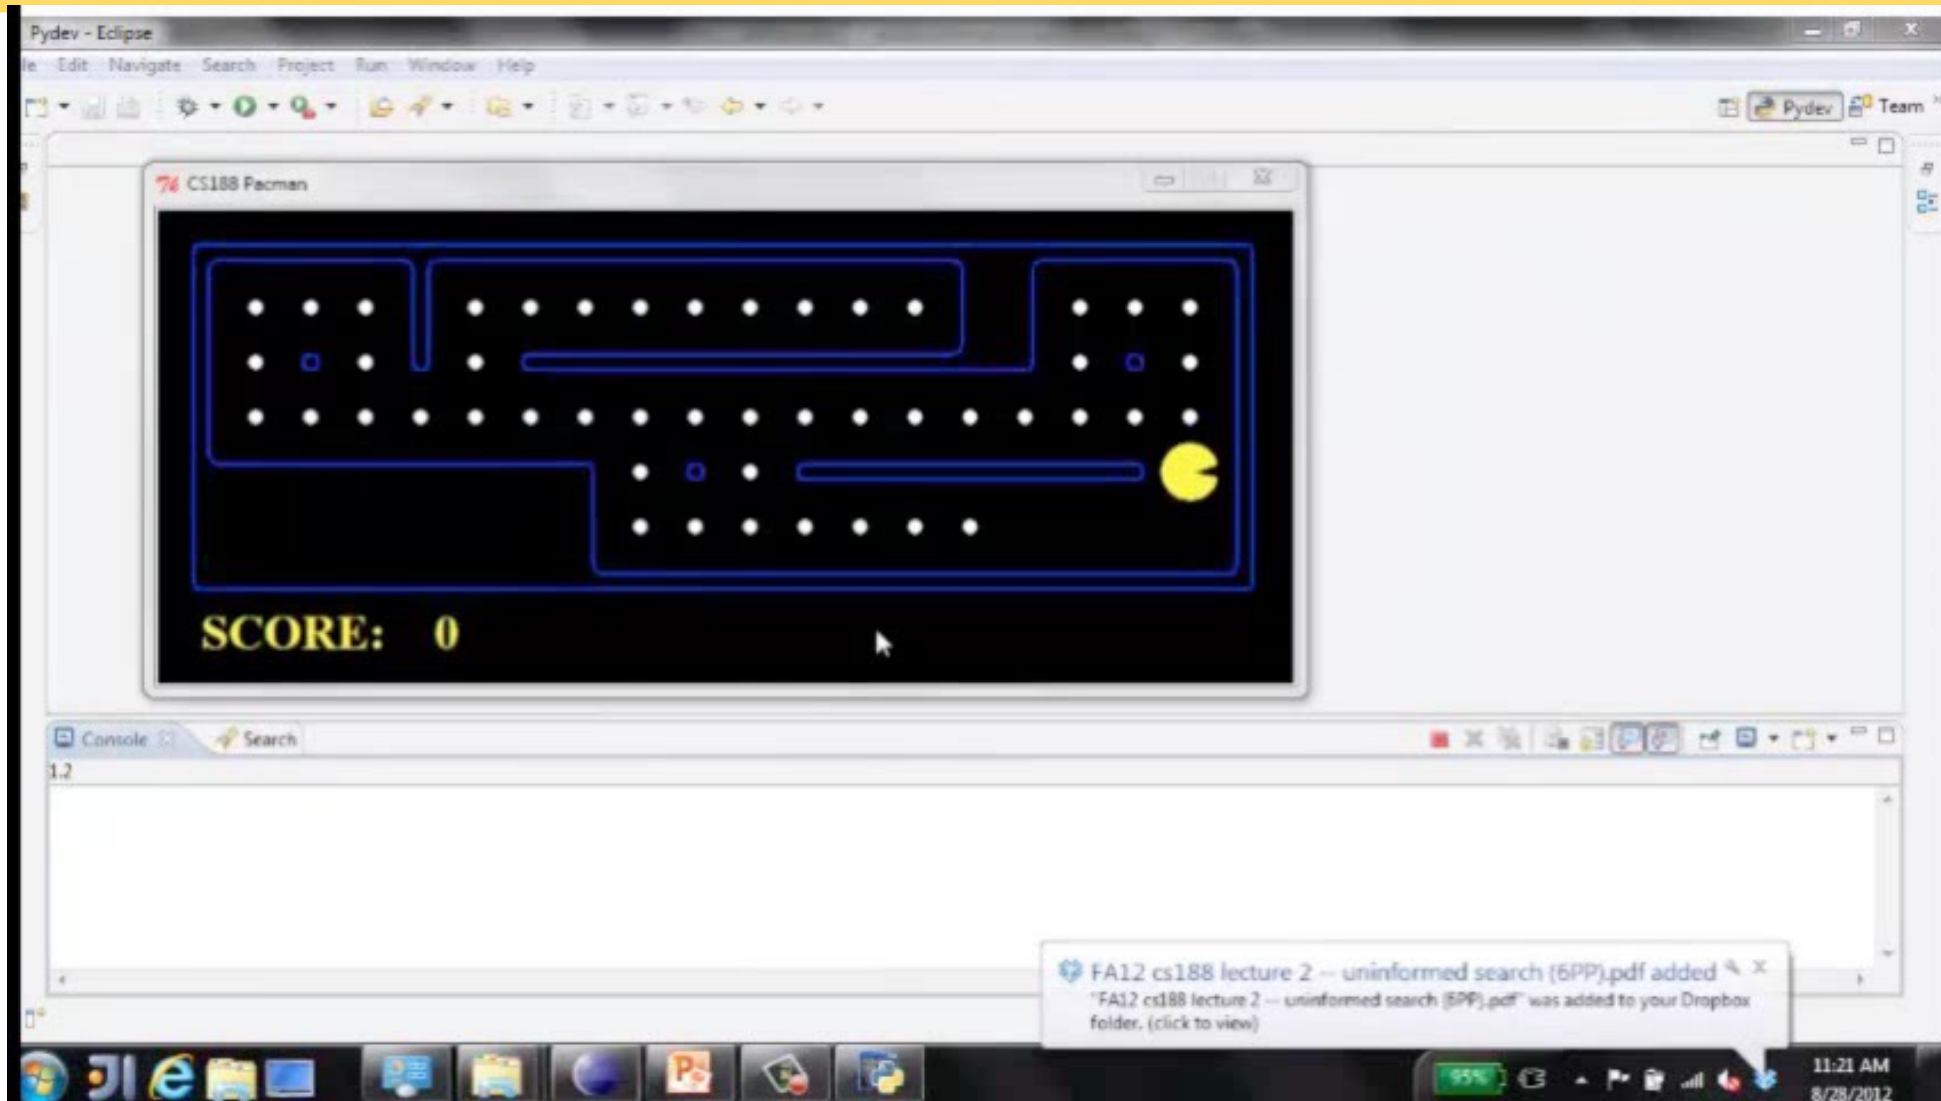
- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - Do not consider future consequences of their actions
 - **Consider how the world IS**
- Can a reflex agent be rational?



Video of Demo Reflex Optimal

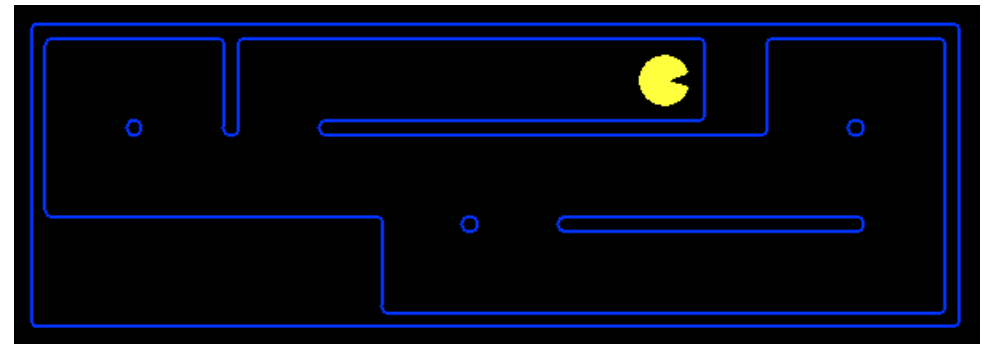
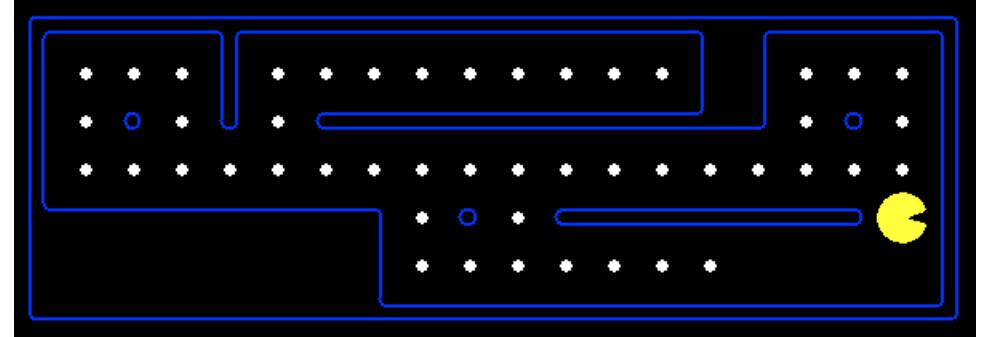


Video of Demo Reflex Odd

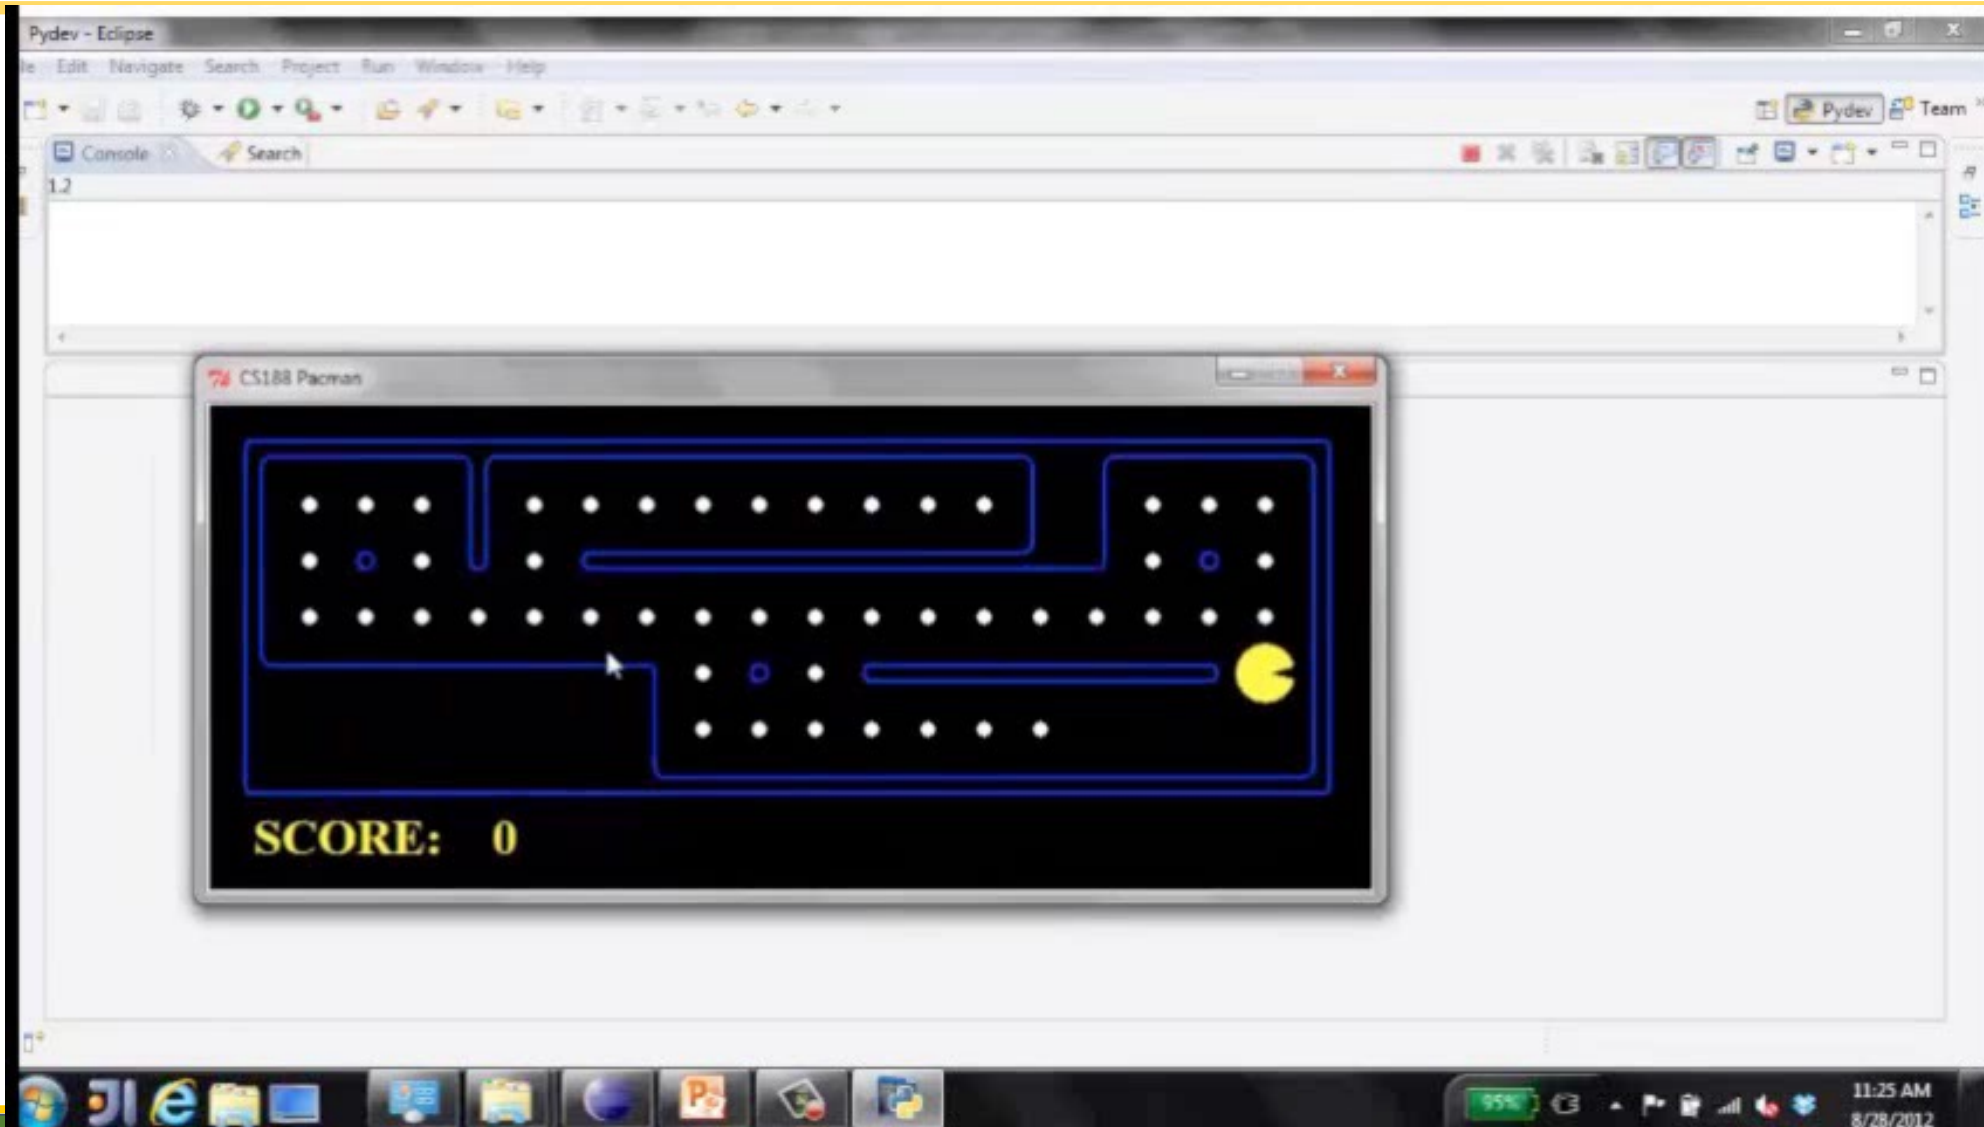


Goal-based Agents

- Goal-based agents:
 - Plan ahead
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - **Act on how the world WOULD BE**



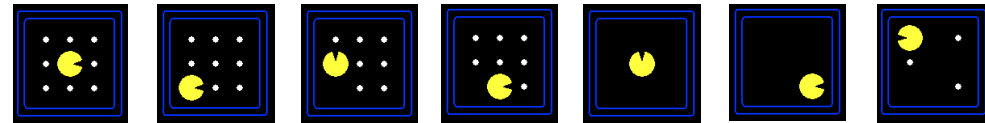
Video of Demo Mastermind



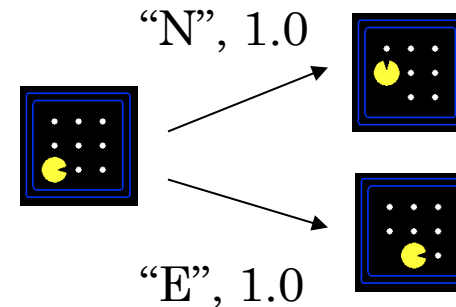
Search Problem

- A **search problem** consists of:

- A state space



- A successor function
(with actions, costs)



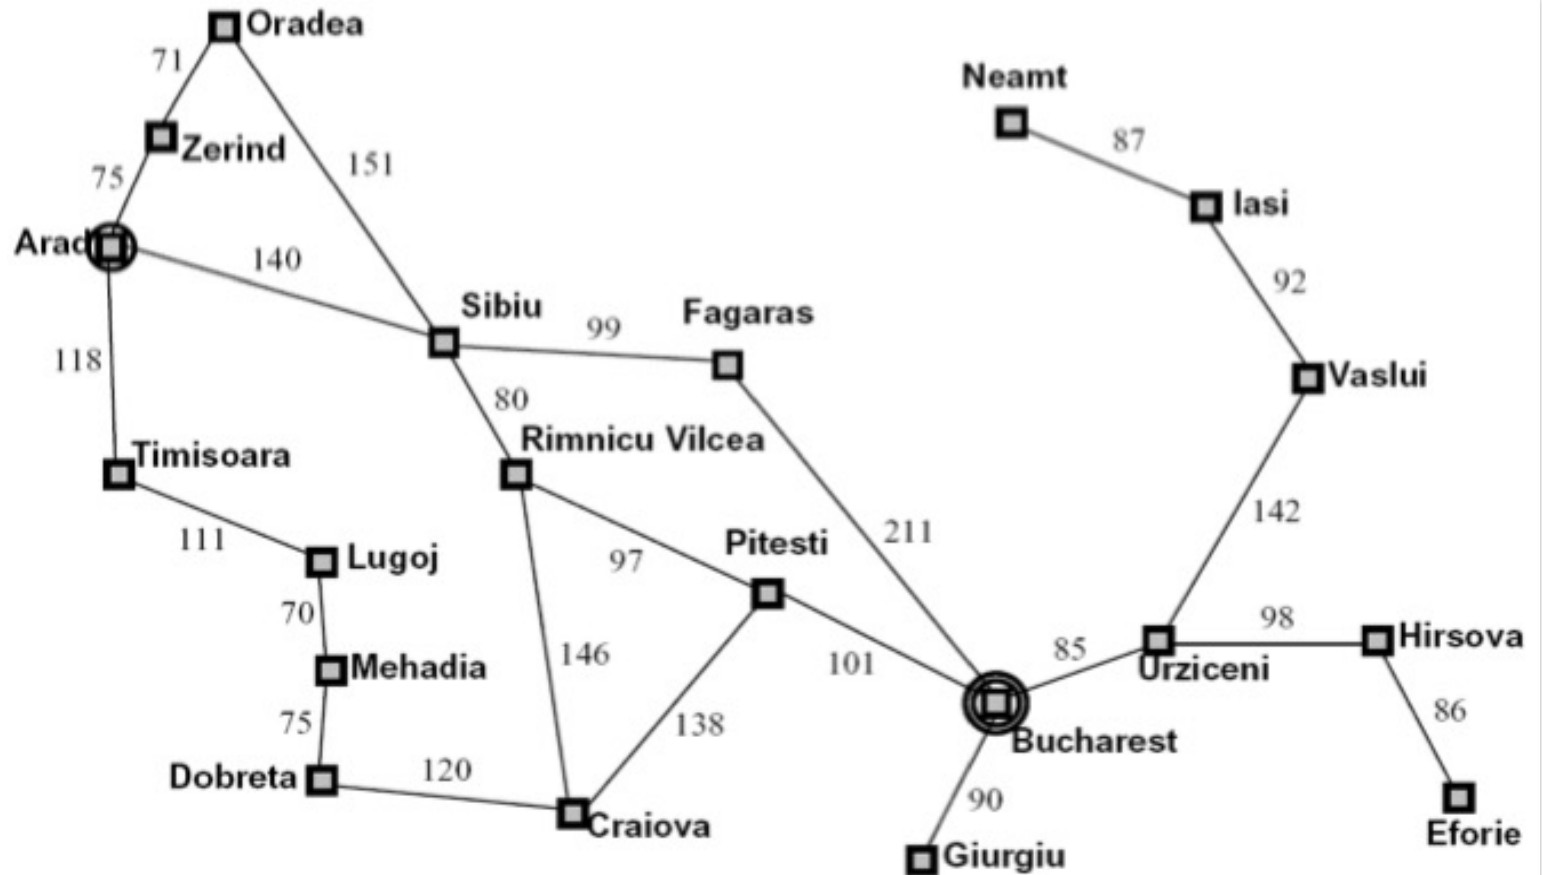
- A start state and a goal test

- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state



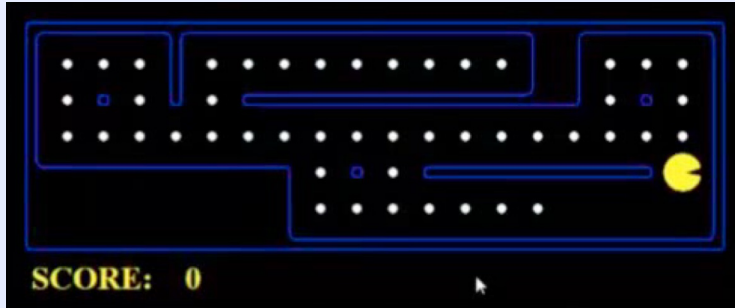
Example: Romania

- State space:
 - Cities
- Successor function:
 - Go to adj city with cost = dist
- Start state:
 - Arad
- Goal test:
 - Is state == Bucharest?
- Solution?



What is in State Space

The **world state** includes every last detail of the environment

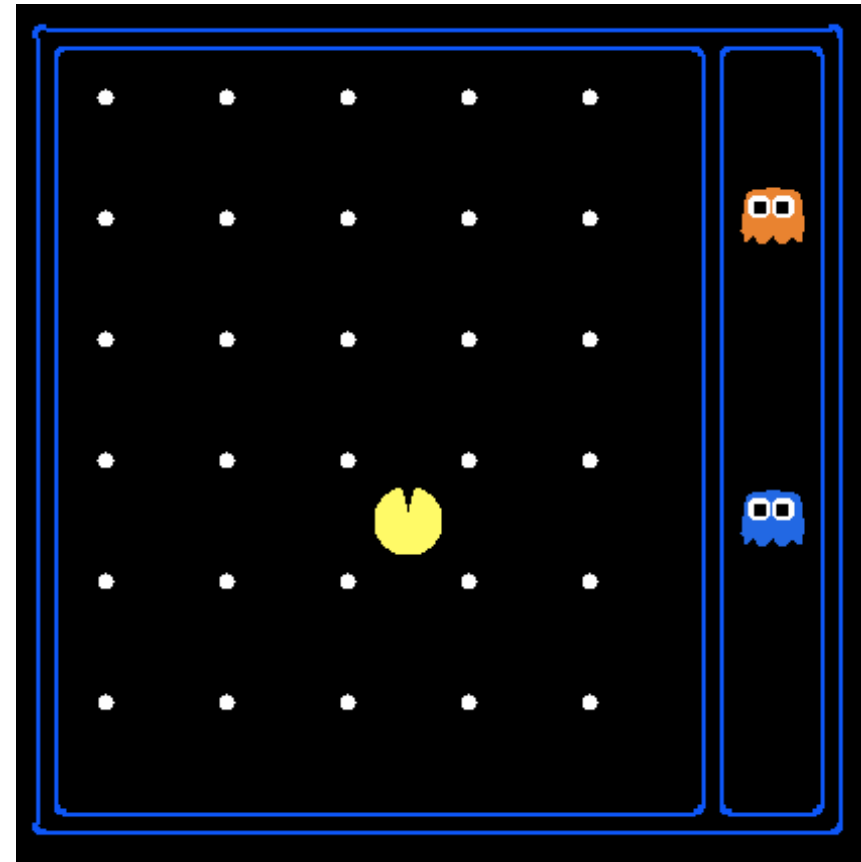


- Problem: Pathing
 - States: (x,y) location
 - Actions: NSEW
 - Successor: update location only
 - Goal test: is $(x,y)=\text{END}$

- Problem: Eat-All-Dots
 - States: $\{(x,y), \text{dot booleans}\}$
 - Actions: NSEW
 - Successor: update location and possibly a dot boolean
 - Goal test: dots all false

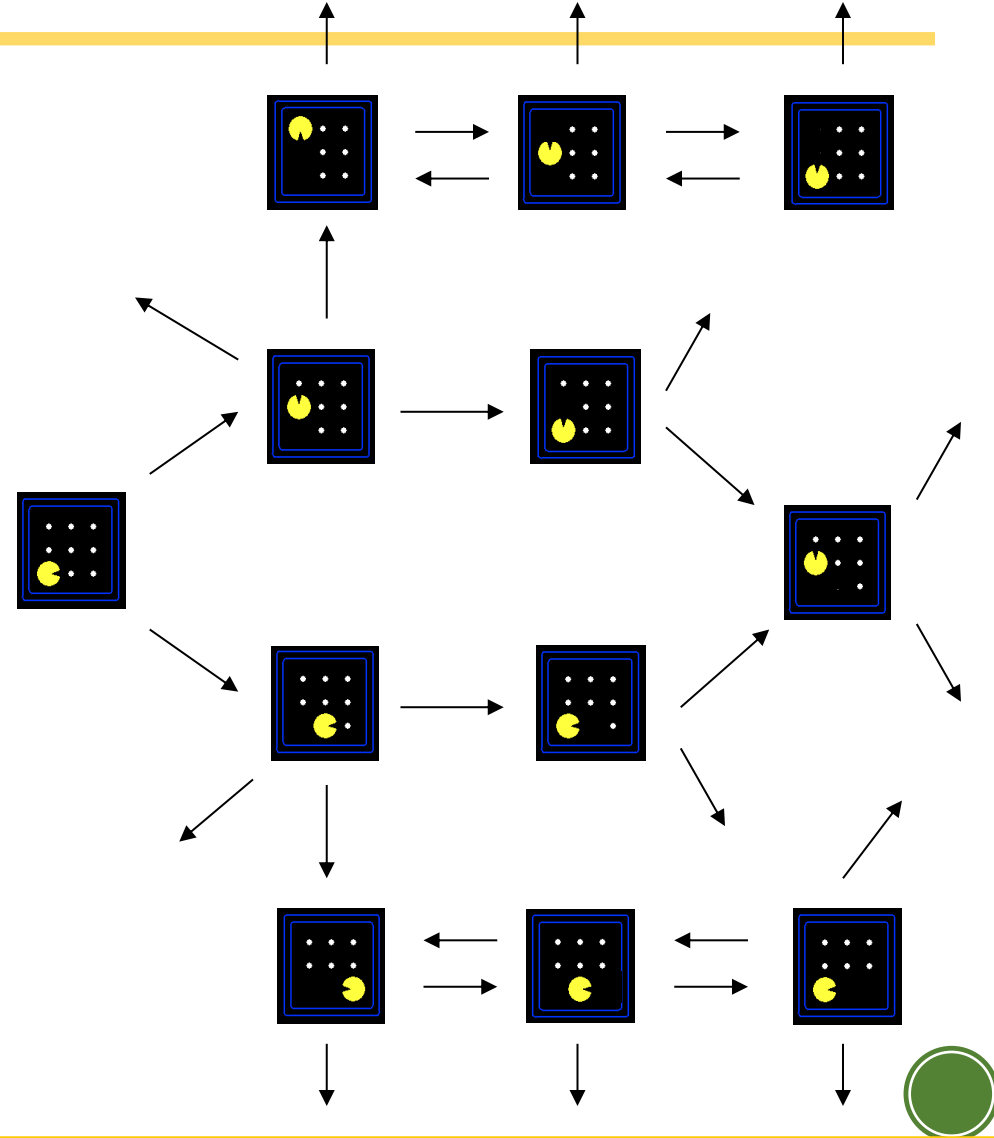
State Space Size

- Search Problem:
Eat all of the food
- Pacman positions: $10 \times 12 = 120$
- Pacman facing: up, down, left, right
- Food Count: 30
- Ghost positions: 12
- How many
- World states? $120 \cdot (2^{30}) \cdot (12^2) \cdot 4$
- States for pathing? 120
- States for eat-all-dots? $120 \cdot (2^{30})$



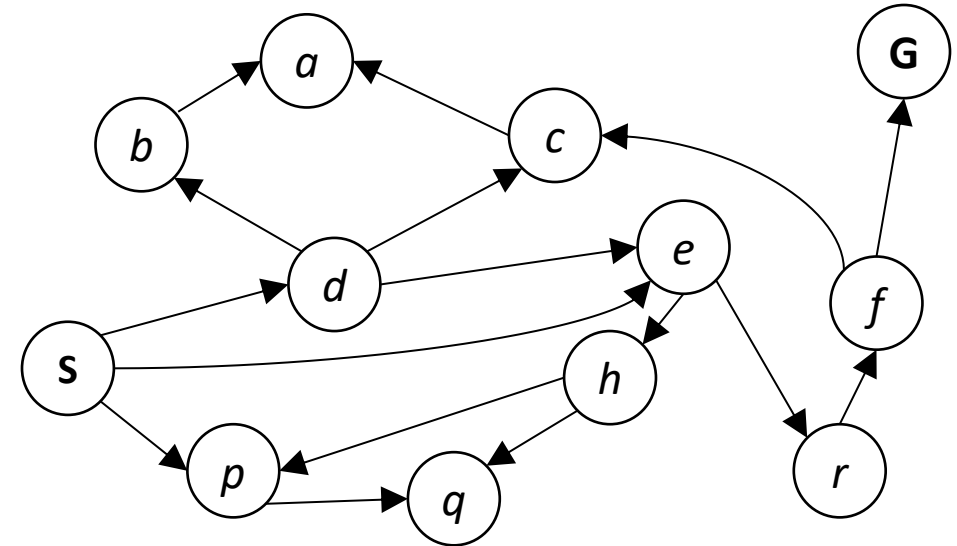
State Space Graphs

- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



State Space Graphs

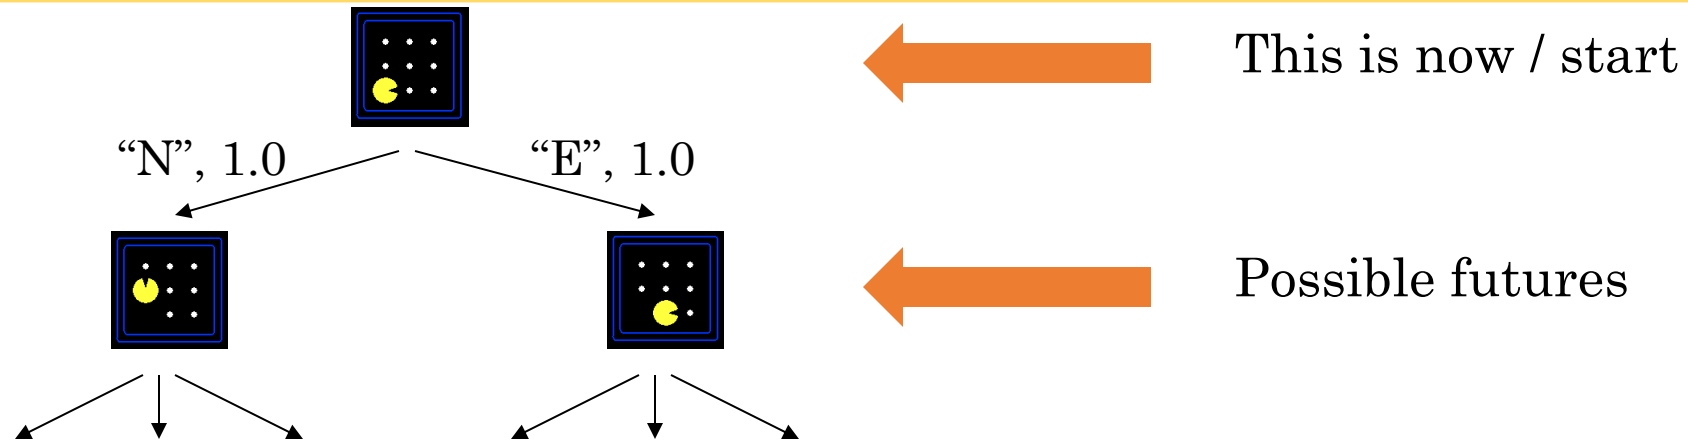
- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



Tiny state space graph for a tiny search problem



Search Trees

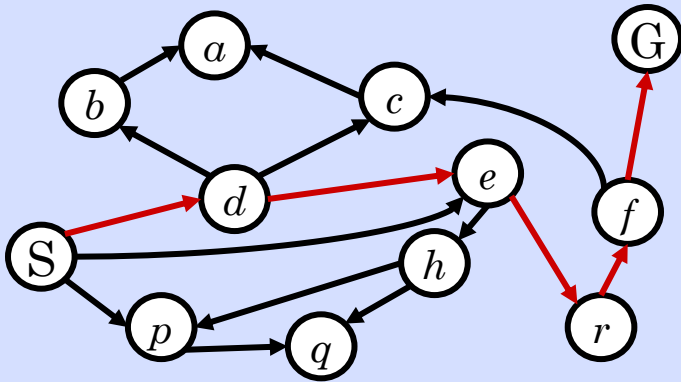


- A search tree:
 - A “what if” tree of plans and their outcomes
 - The start state is the root node
 - Children correspond to successors
 - Nodes show states, but correspond to PLANS that achieve those states
 - For most problems, we can never actually build the whole tree



State Space Graphs vs. Search Trees

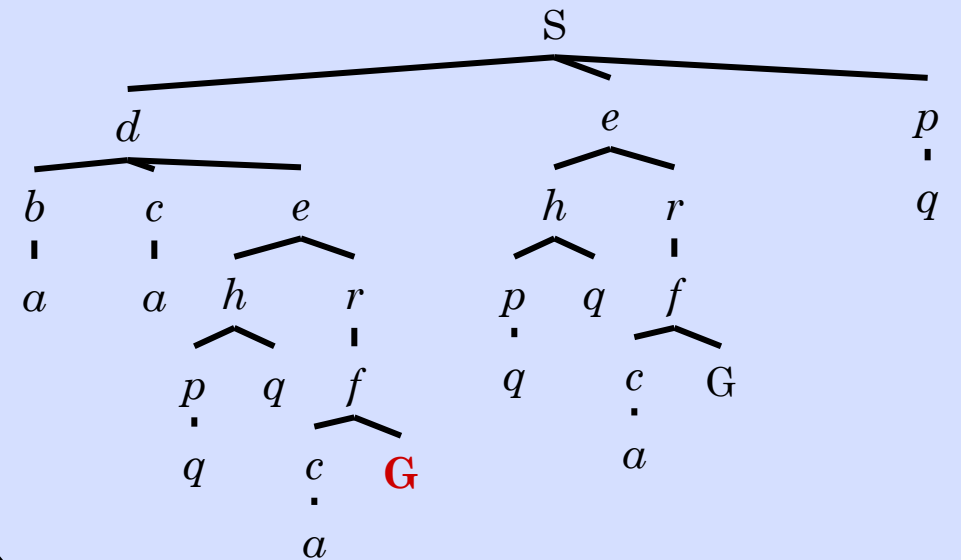
State Space Graph



*Each NODE in
in the search tree
is an entire
PATH in the
state space
graph.*

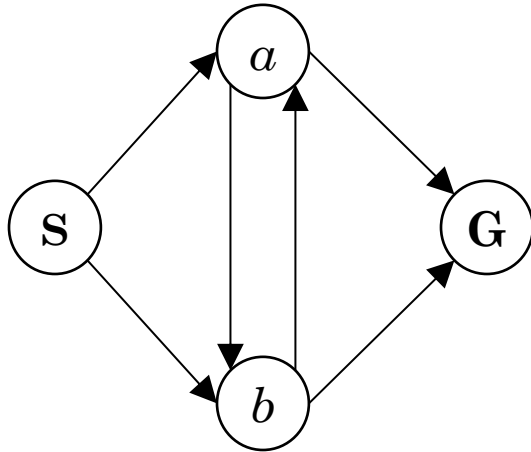
*We construct both
on demand – and
we construct as
little as possible.*

Search Tree



Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:



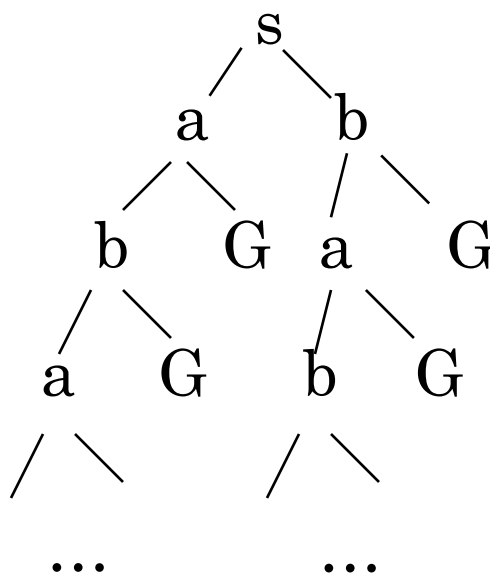
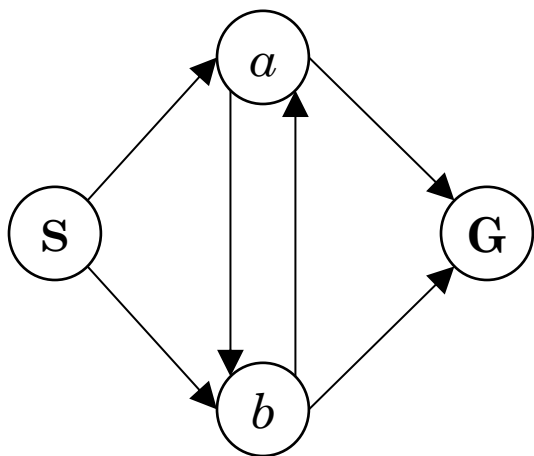
How big is its search tree (from S)?



Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:

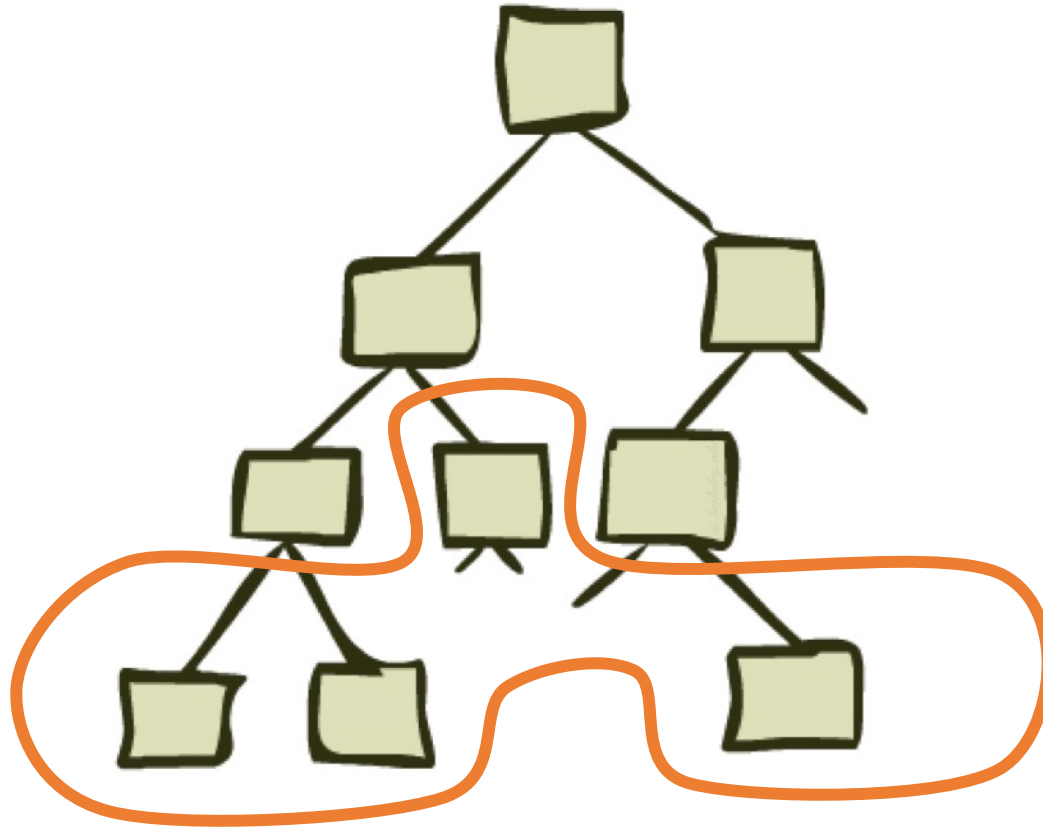
How big is its search tree (from S)?



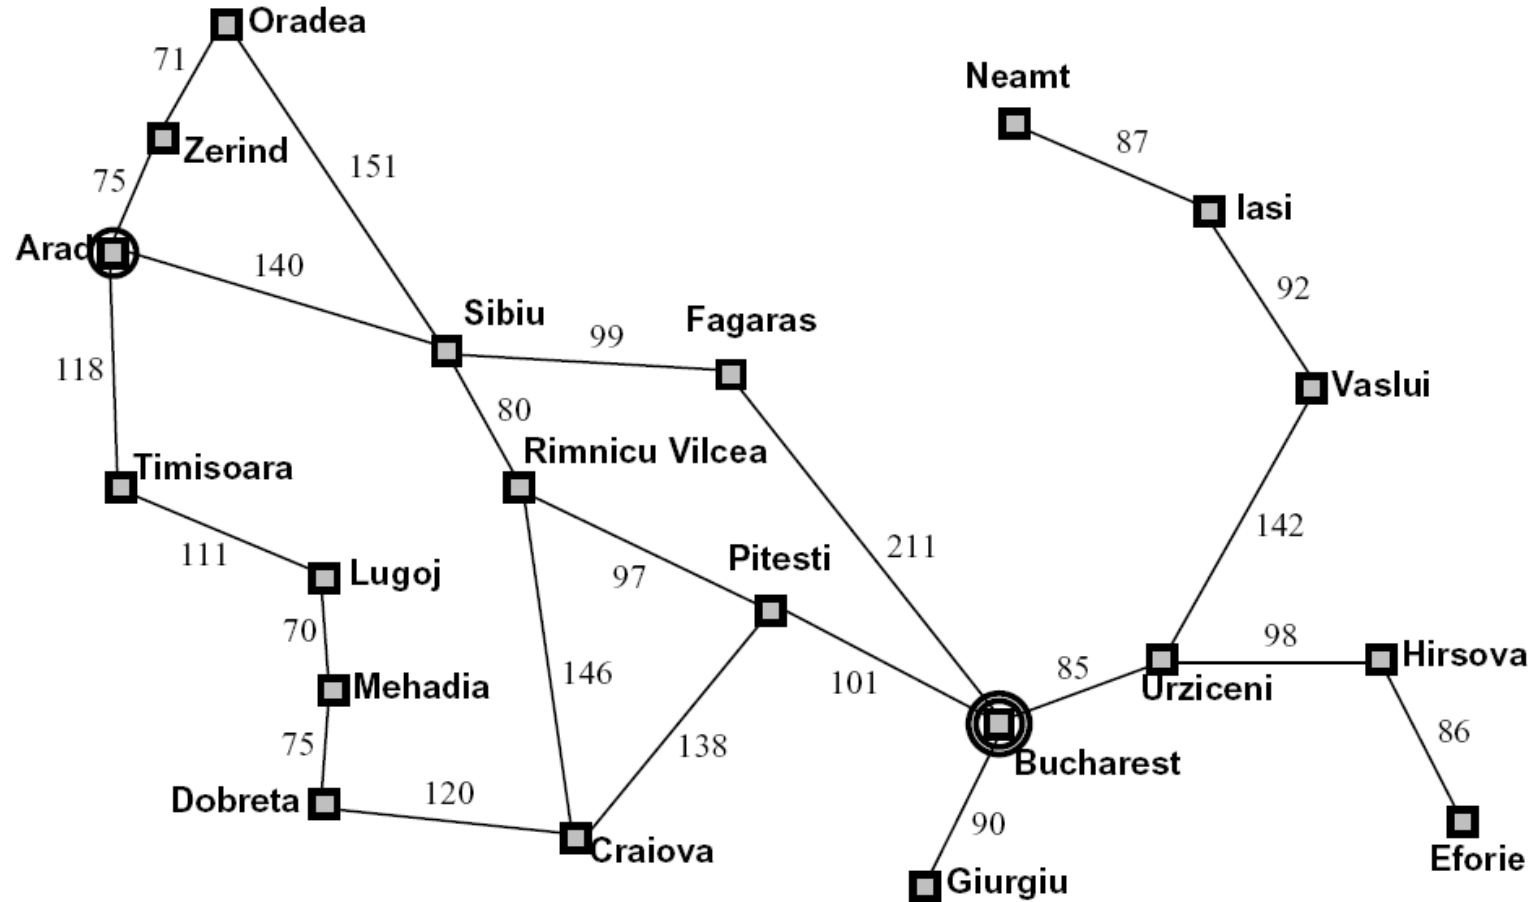
Important: Lots of repeated structure in the search tree!



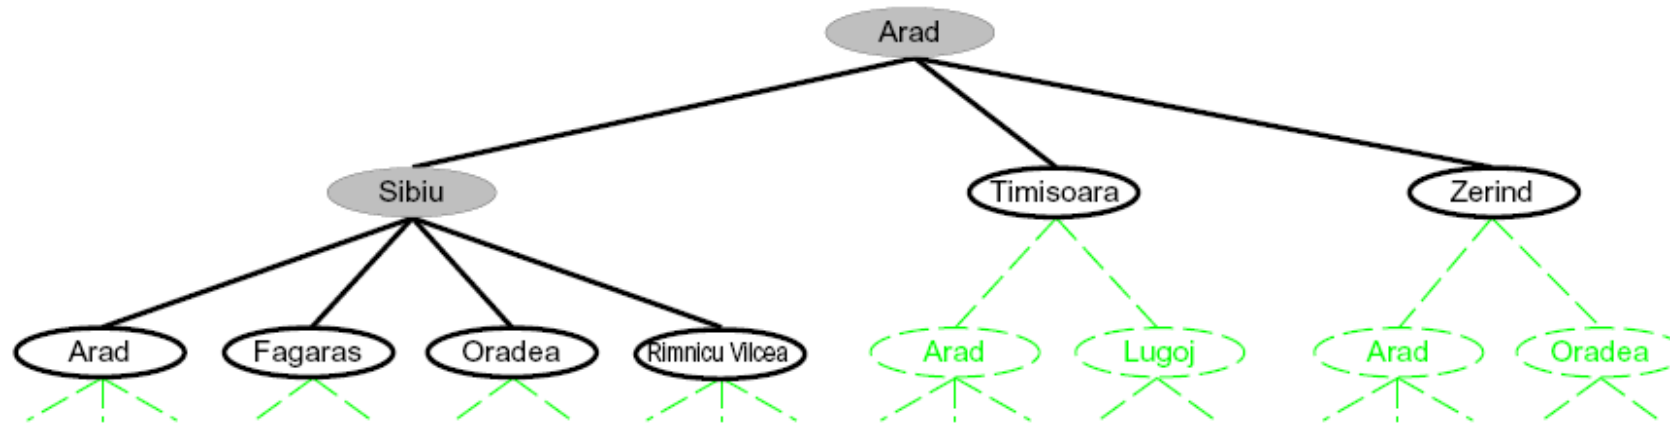
Tree Search



Search Example: Romania



Searching with a Search Tree



- Search:
 - Expand out potential plans (tree nodes)
 - Maintain a **fringe** of partial plans under consideration
 - Try to expand as few tree nodes as possible

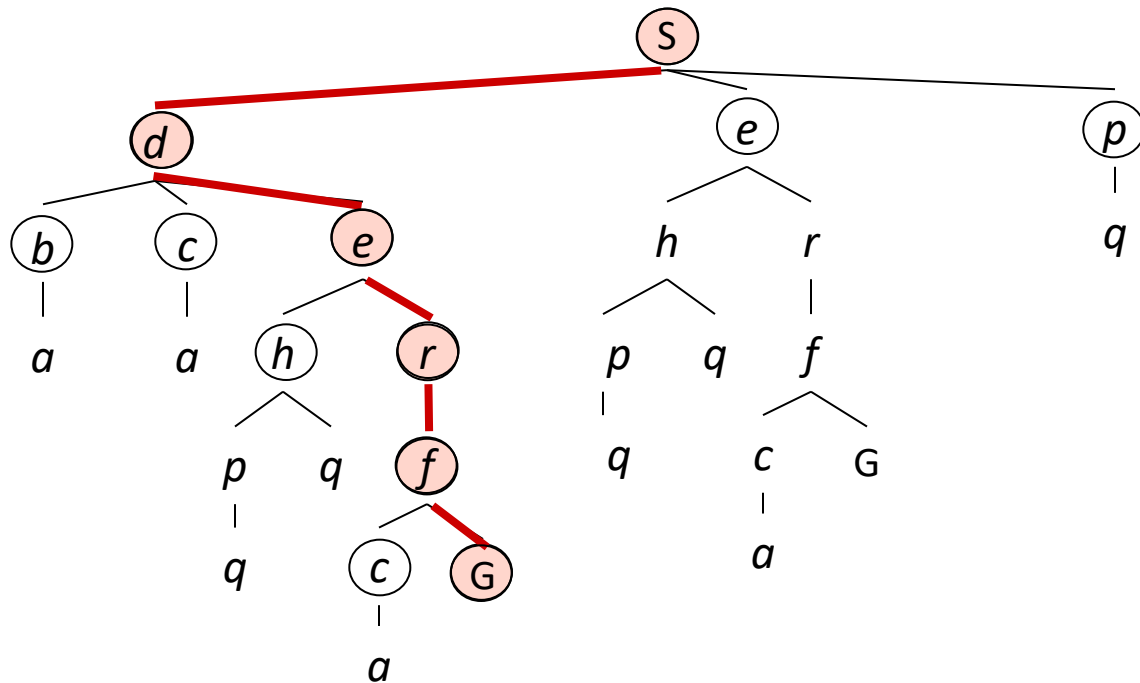
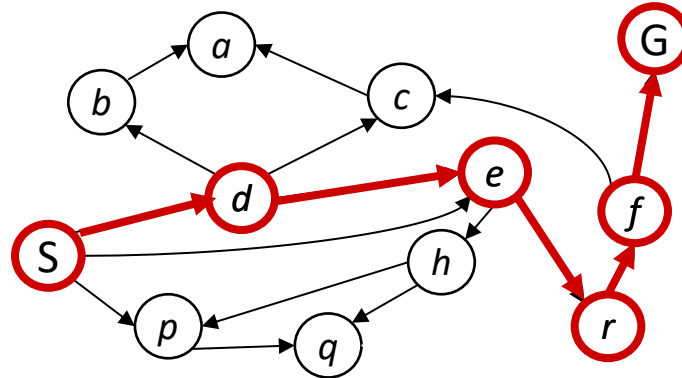


General Tree Search

- Tree Search
 - Initialize the *root node* of the search tree with the *start state*
 - While there are unexpanded leaf nodes (fringe):
 - Choose a leaf node (strategy)
 - If the node contains a goal state:
return the corresponding solution
 - Else: expand the node and add its children to the tree
- Important ideas:
 - Fringe
 - Expansion
- Strategy: which fringe nodes to explore?



Example: Tree Search



~~s~~
~~s → d~~
s → e
s → p
s → d → b
s → d → c
~~s → d → e~~
s → d → e → h
~~s → d → e → r~~
~~s → d → e → r → f~~
s → d → e → r → f → c
~~s → d → e → r → f → G~~

