
CS 471/571 (Fall 2023): Introduction to Artificial Intelligence

Lecture 25: Neural Nets

Thanh H. Nguyen

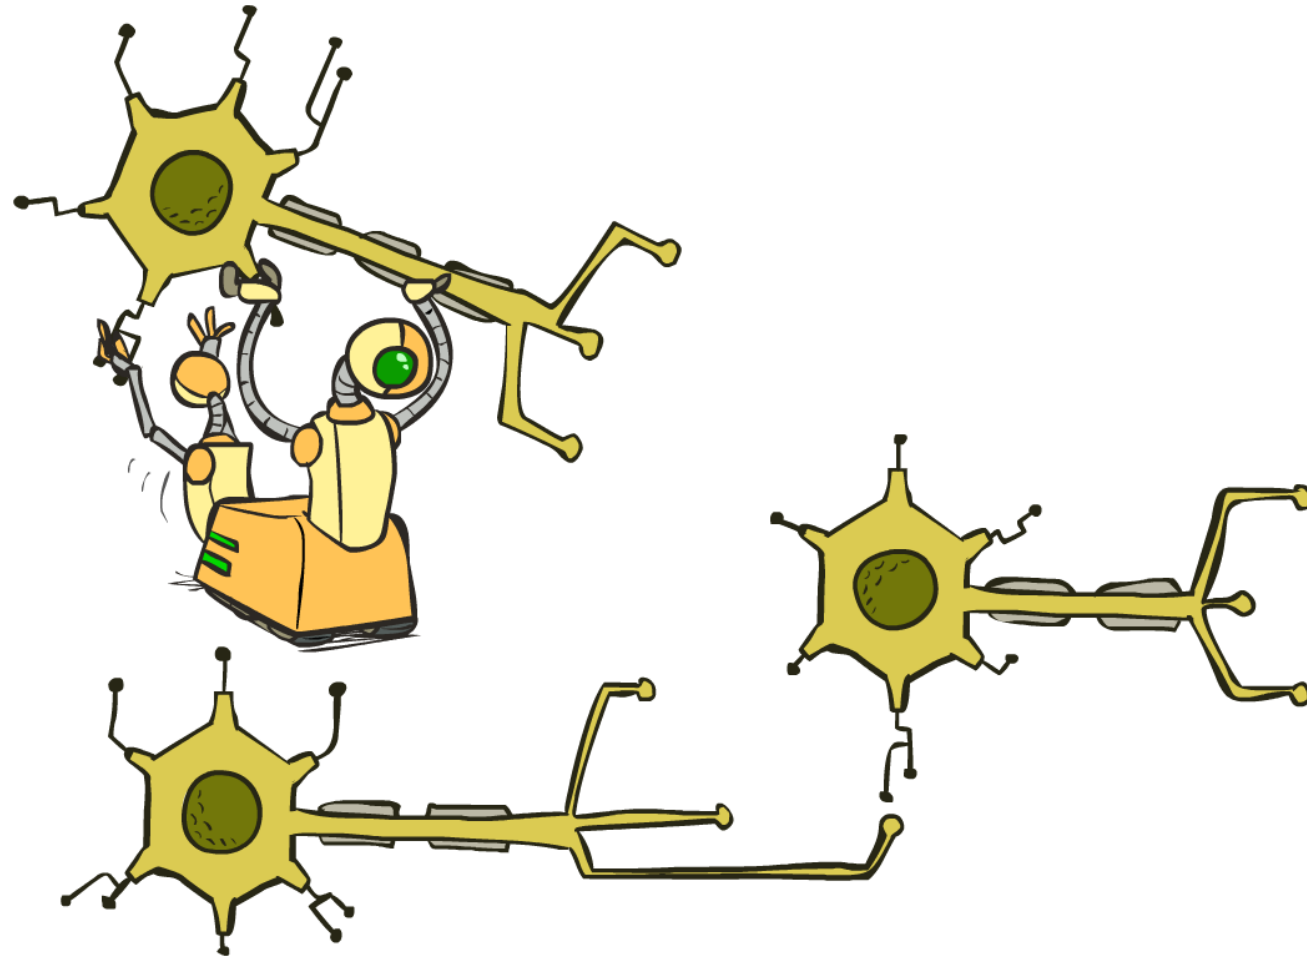
Source: <http://ai.berkeley.edu/home.html>



Announcement & Reminder

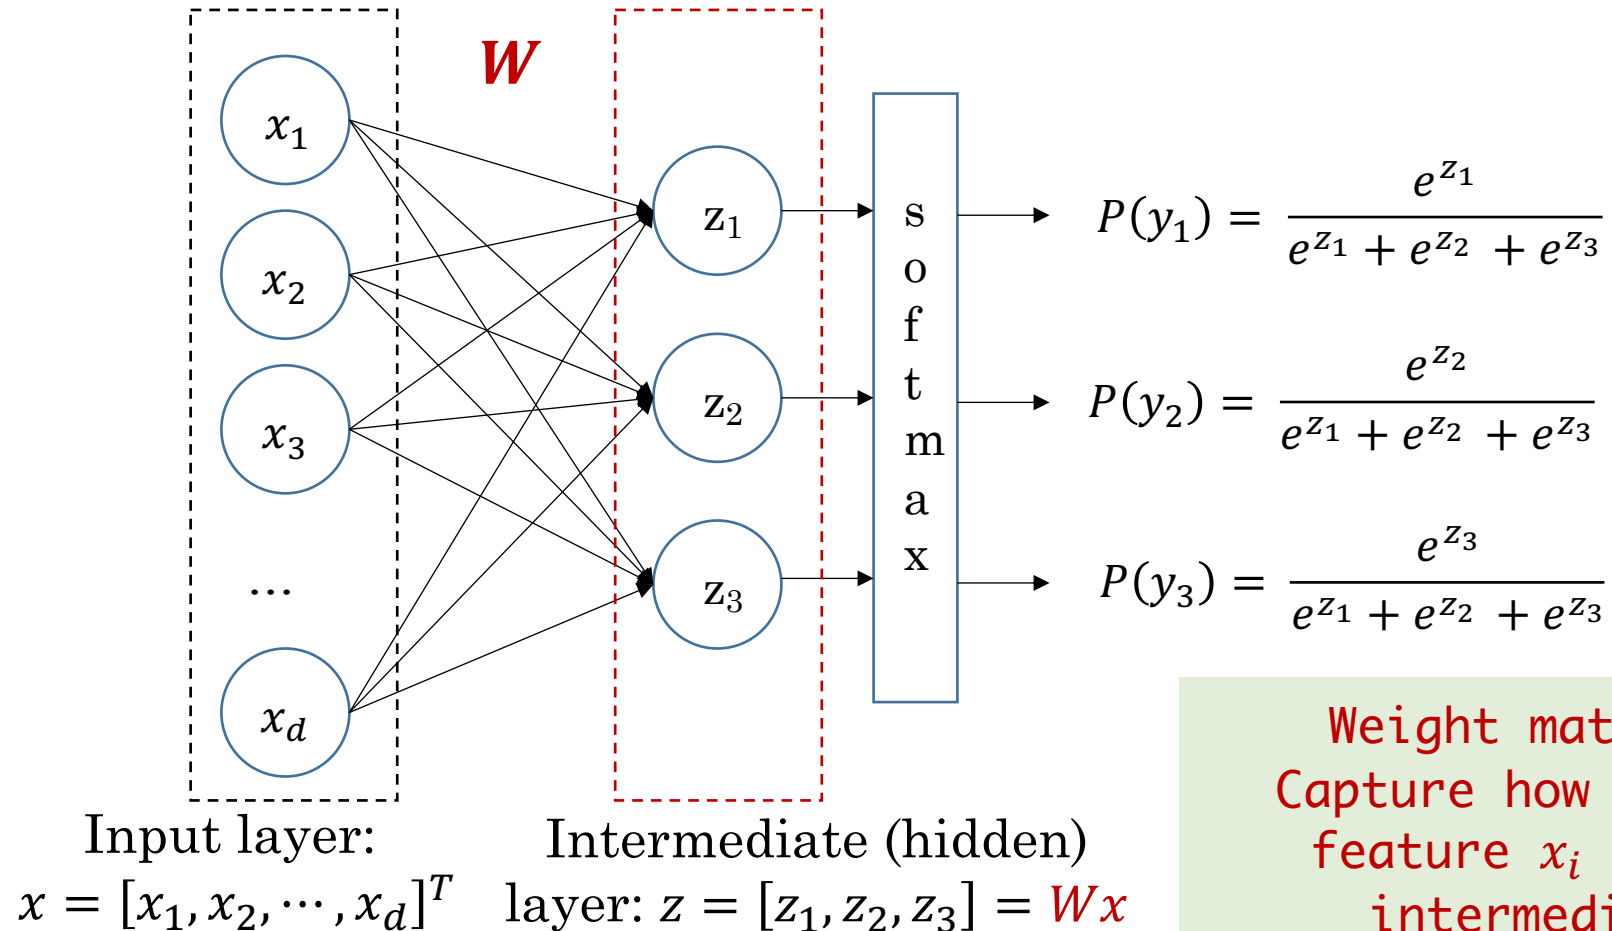
- Exam Review
 - Friday, Dec 01st, 2023
- Written assignment 4
 - Deadline: Wednesday, November 29th, 2023.
- Student experience survey
 - Deadline: 06:00 AM on Monday, Dec 4th, 2023
 - If $\geq 80\%$ of students complete the survey, everyone will get an extra 2% credit for your final grade

Neural Networks



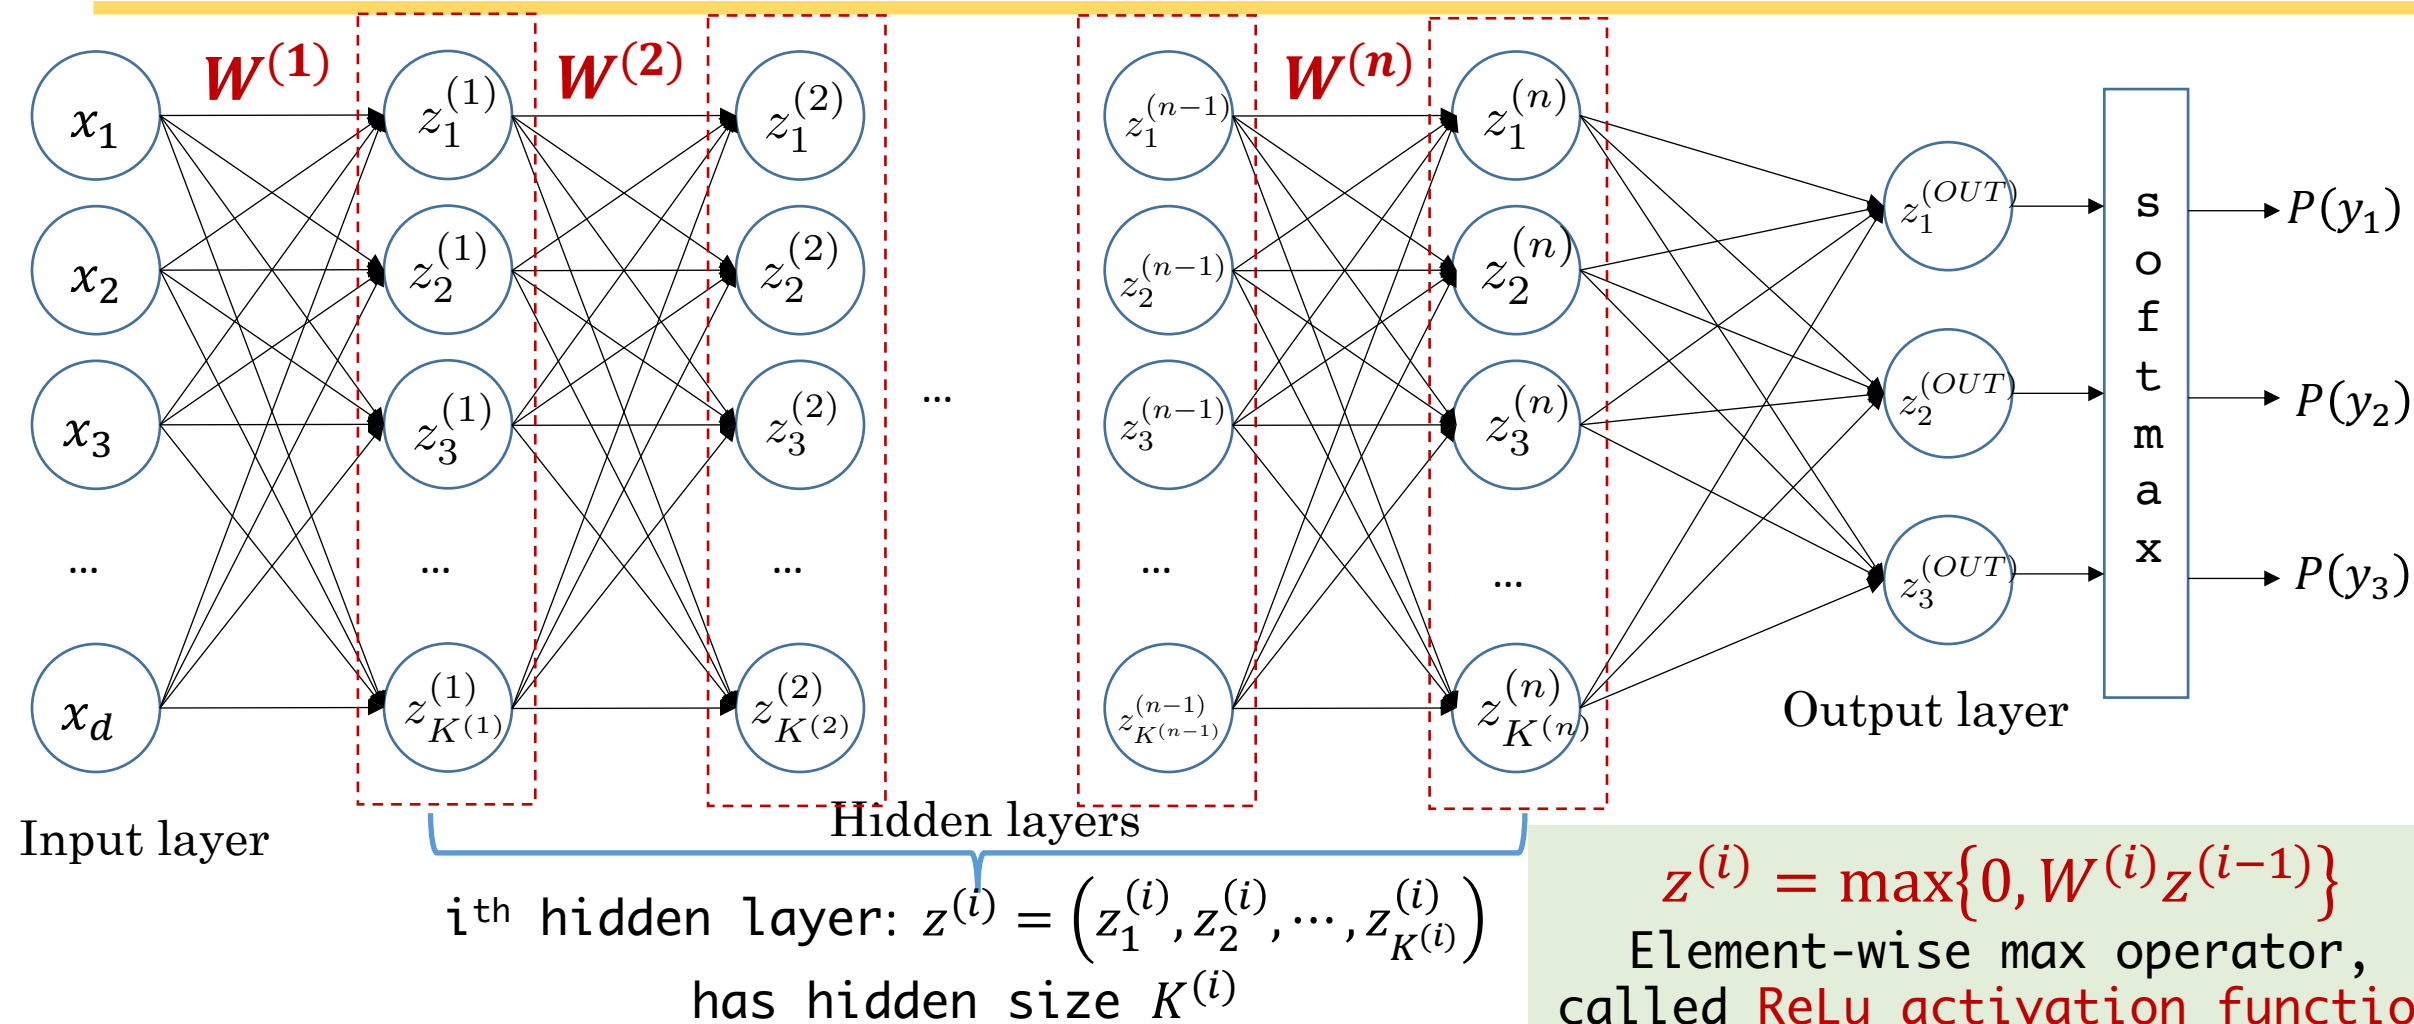
Multi-class Logistic Regression

- = special case of neural network



Weight matrix: $W = \{w_{ij}\}$
Capture how much each input feature x_i influence each intermediate value z_j

Deep Neural Network = Learn the Features!

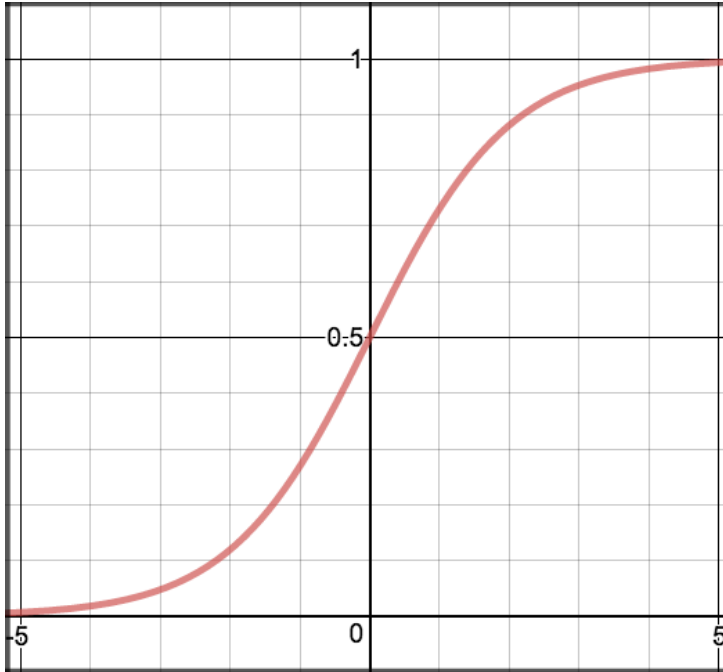


Common Activation Functions

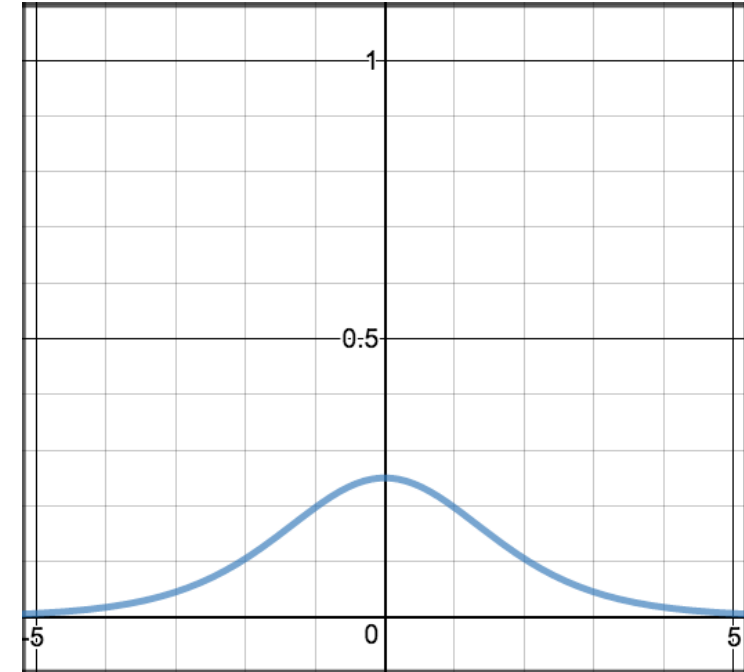
Source: https://ml-cheatsheet.readthedocs.io/en/latest/activation_functions.html



Sigmoid



Function: $z = \frac{1}{1 + e^{-x}}$



Derivative: $\frac{dz}{dx} = z \cdot (1 - z)$

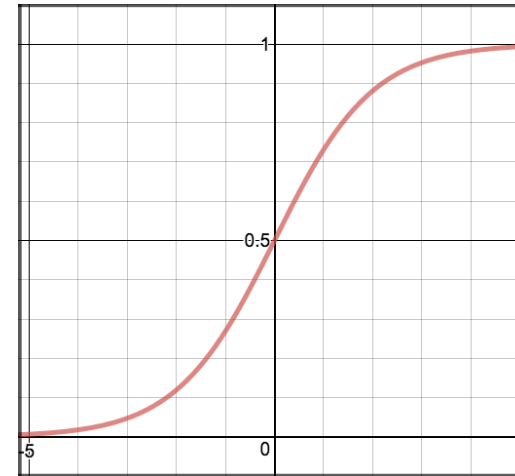
Sigmoid

■ Pros

- Is nonlinear.
- Has a smooth gradient.
- Good for a classifier.
- The output is bounded within (0,1)

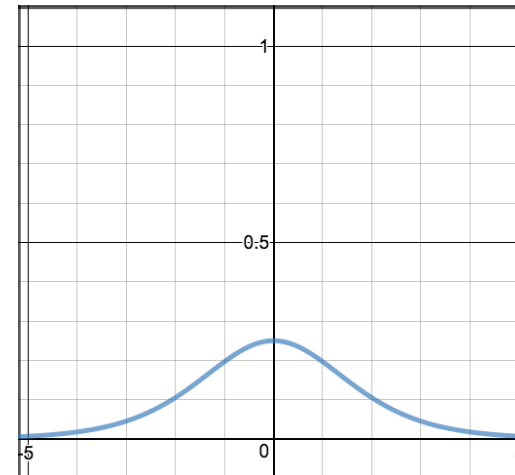
■ Cons

- Towards either end of the sigmoid function, the z values tend to respond very less to changes in x.
- Gives rise to a problem of “vanishing gradients”.



Function:

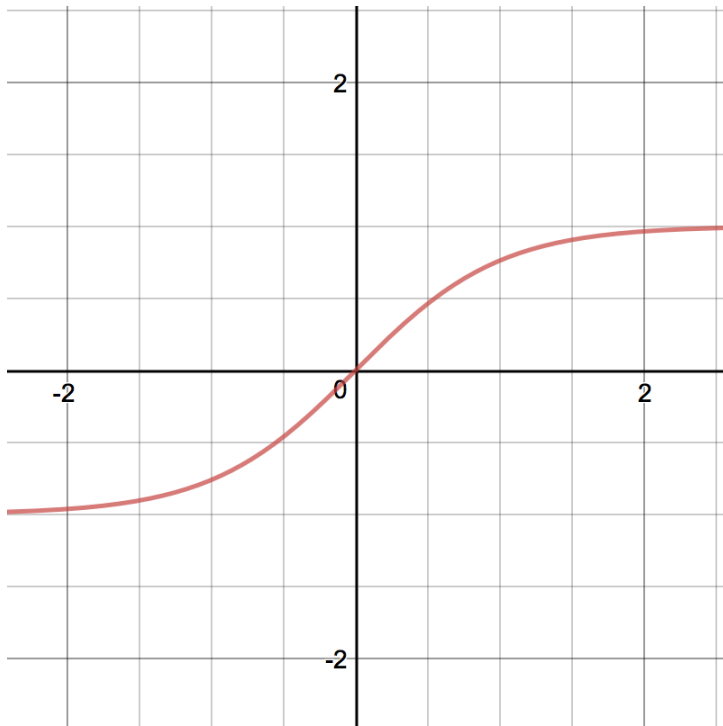
$$z = \frac{1}{1 + e^{-x}}$$



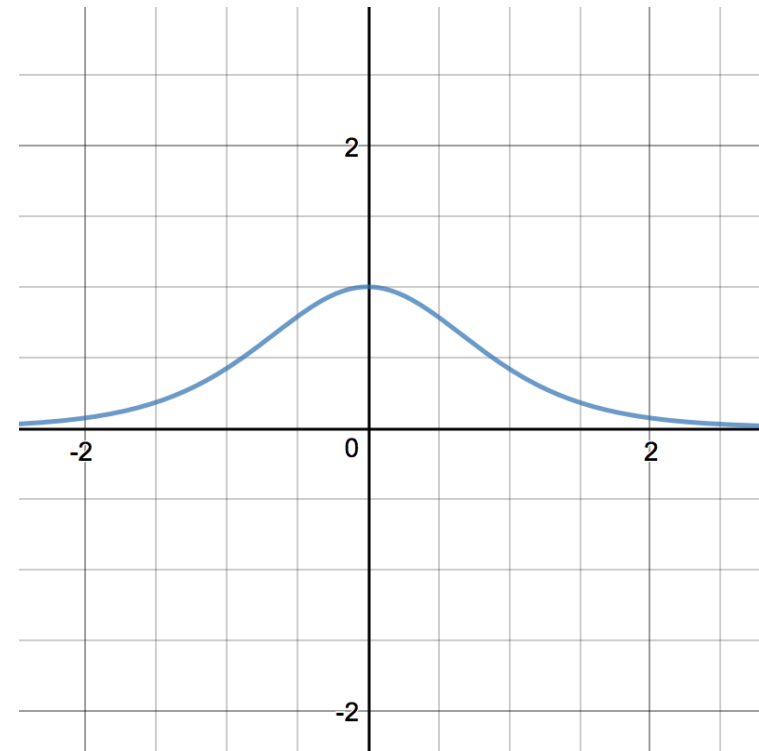
Derivative:

$$\frac{dz}{dx} = z \cdot (1 - z)$$

Tanh



$$\text{Function: } z = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



$$\text{Derivative: } \frac{dz}{dx} = 1 - z^2$$

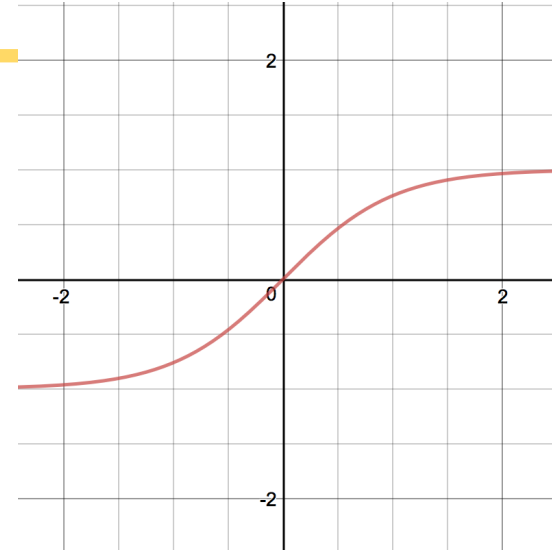
Tanh

- Pros

- The gradient is stronger for tanh than sigmoid (derivatives are steeper).

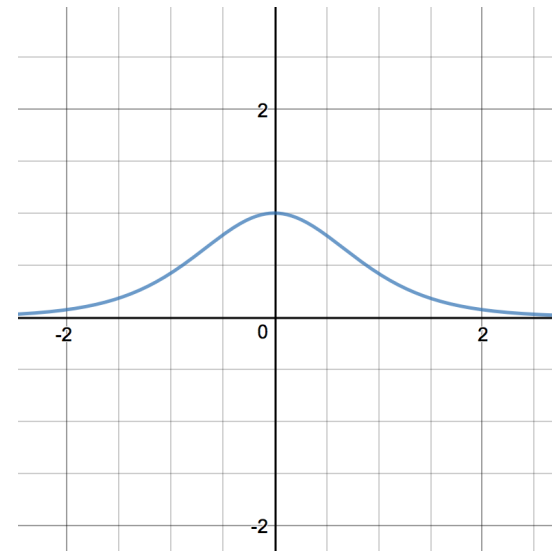
- Cons

- Tanh also has the vanishing gradient problem.



Function:

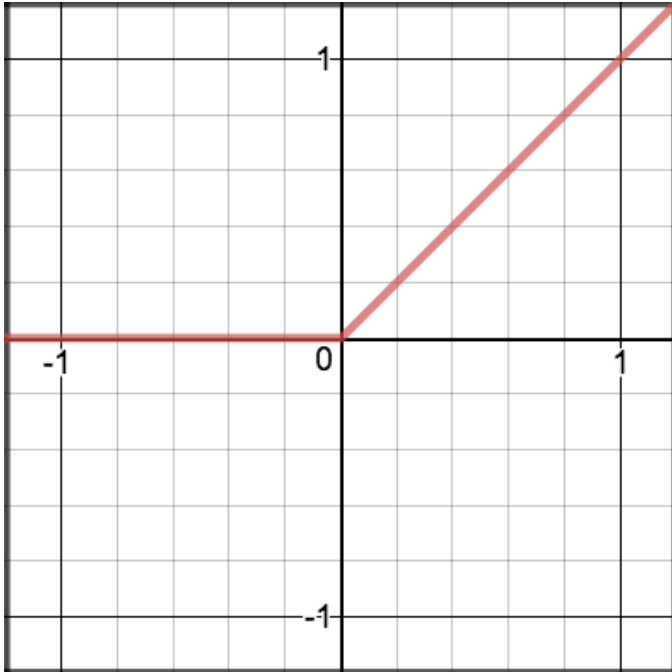
$$z = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



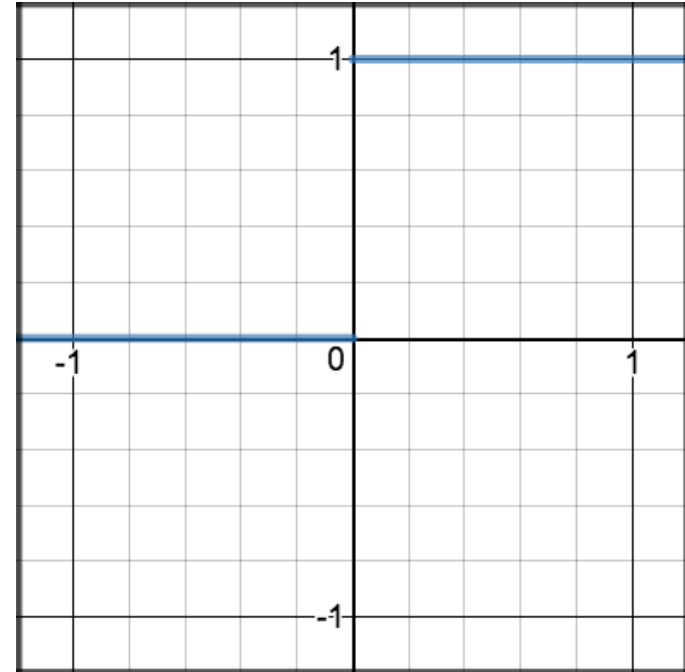
Derivative:

$$\frac{dz}{dx} = 1 - z^2$$

ReLU



Function: $z = \max\{0, x\}$



Derivative: $\frac{dz}{dx} = \begin{cases} 1, & \text{if } x > 0 \\ 0, & \text{if } x < 0 \end{cases}$



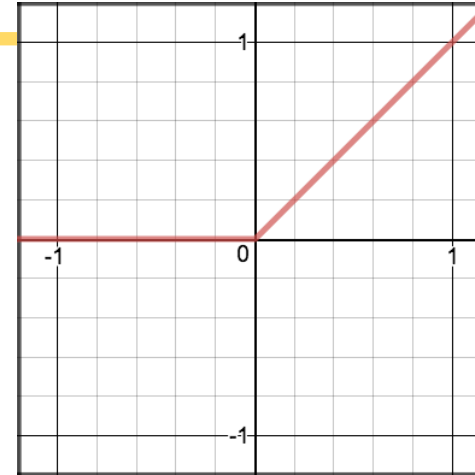
ReLU

■ Pros

- ReLu avoids and rectifies vanishing gradient problem.
- ReLu is less computationally expensive

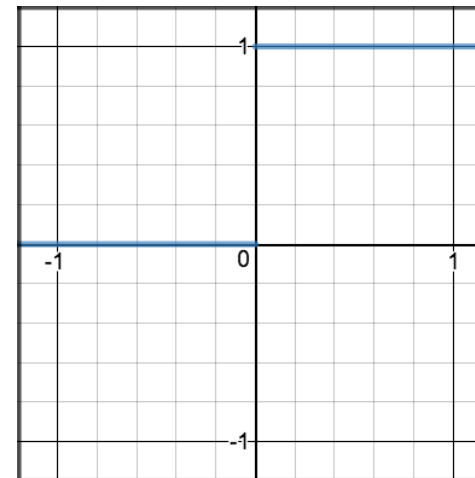
■ Cons

- ReLu should only be used within hidden layers.
- Dying ReLu problem: When $(x < 0)$, gradient of ReLu is 0, meaning the weights will not get adjusted during gradient descent.
- The range of ReLu is $[0, \infty)$. This means it can blow up the activation.



Function:

$$z = \max\{0, x\}$$

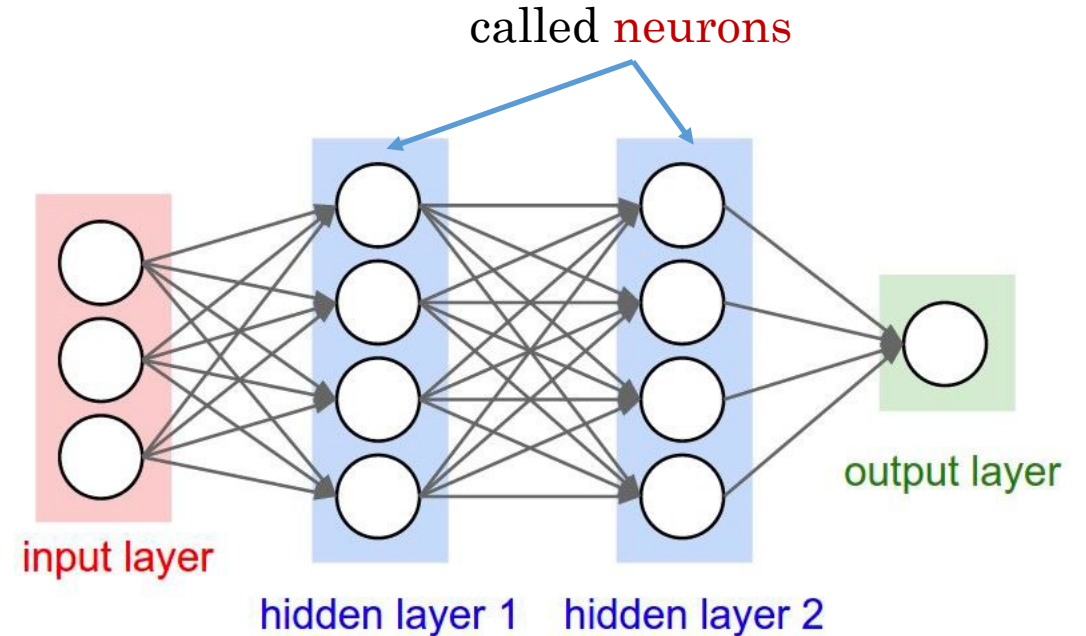


Derivative:

$$\frac{dz}{dx} = \begin{cases} 1, & \text{if } x > 0 \\ 0, & \text{if } x < 0 \end{cases}$$

Train a Neural Net

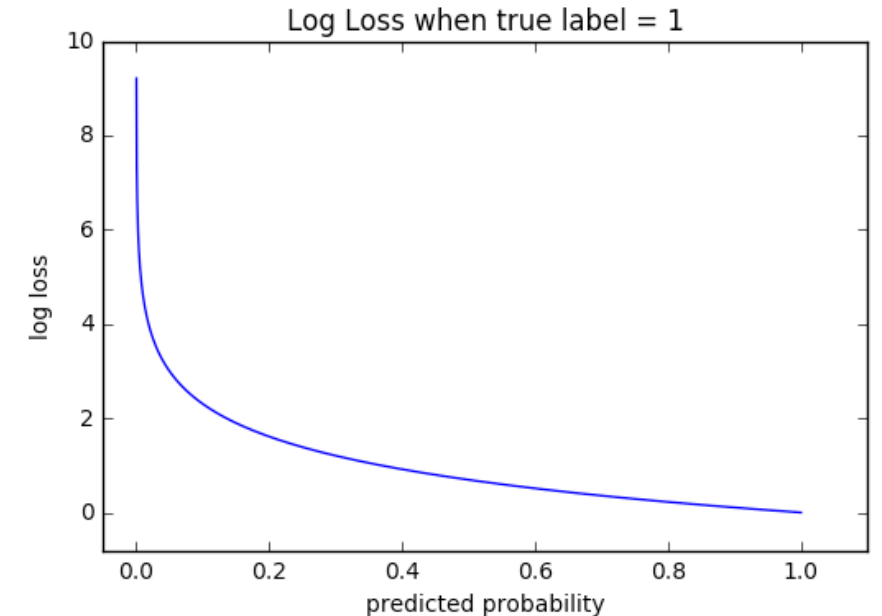
- Train a neural networks: gradient descent
 - **Forward pass**: input is passed forward through the network to produce a prediction output. A loss metric is computed based on the difference between prediction and target (true output).
 - **Backward pass**: derivatives of this loss metric are calculated and propagated back through the network using a technique called backpropagation.
- We will talk more about this!!!



Common Loss Functions

Cross Entropy

- Measures the performance of a classification model whose output is a probability value between 0 and 1
- Cross-entropy loss increases as the predicted probability diverges from the actual label.
- A perfect model would have a log loss of 0.



- Binary classification: $Loss = -(y \log p + (1 - y) \log(1 - p))$
- Multi-classification: $Loss = -\sum_l y_l \log p_l$
 - $y_l = 1$ if the true label is l , and $y_l = 0$, otherwise
 - p_l : predicted probability of having label l

Mean Square Error (MSE)

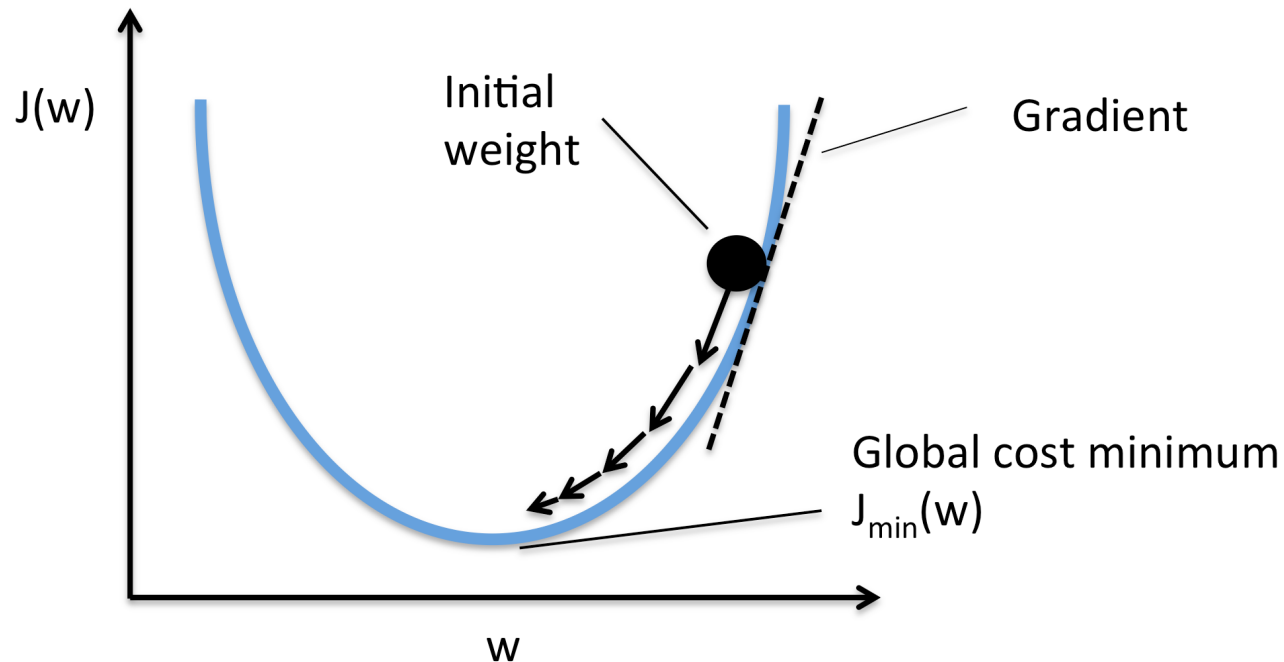
- Commonly used for regression loss
- Measure the average of squared difference between predictions and actual observations

$$MSE = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$

Training a Neural Network

Minimizing Error: Gradient Descent

- Main technique to minimize prediction error in neural network



- Challenge: computing gradient is non-trivial!!!

(Bad Choice) Direct Computation

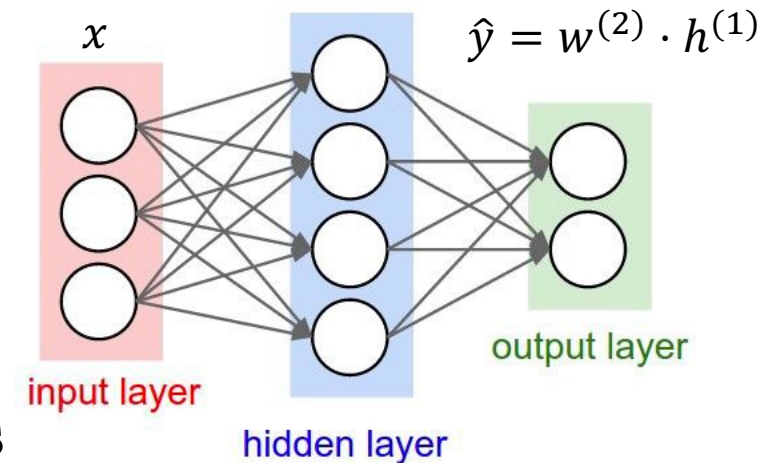
- Input:

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_d \end{bmatrix}$$

- One hidden layer $h^{(1)} = \max(0, w^{(1)}x)$ with $K^{(1)}$ neurons
- Output layer $\hat{y} = w^{(2)} \cdot h^{(1)}$

$$w^{(1)} = \begin{bmatrix} w_{1,1}^{(1)} & \dots & w_{1,K^{(1)}}^{(1)} \\ \vdots & \ddots & \vdots \\ w_{d,1}^{(1)} & \dots & w_{d,K^{(1)}}^{(1)} \end{bmatrix}^T$$

$$w^{(2)} = \begin{bmatrix} w_1^{(2)} \\ w_2^{(2)} \\ \dots \\ w_{K^{(1)}}^{(2)} \end{bmatrix}$$



$$h^{(1)} = \max(0, w^{(1)}x)$$

Hing loss: $\max(0, 1 - \hat{y} \cdot y)$

(Bad Choice) Direct Computation

- Input data: $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})$
- Prediction (two-layer network): $\hat{y}^{(i)} = w^{(2)} \max(0, w^{(1)} x^{(i)})$
- Total loss = prediction loss + regularization

$$\begin{aligned} L &= \sum_i \max(0, 1 - \hat{y}^{(i)} \cdot y^{(i)}) + \lambda R(w^{(1)}) + \lambda R(w^{(2)}) \\ &= \sum_i \max(0, 1 - \hat{y}^{(i)} w^{(2)} \max(0, w^{(1)} x^{(i)})) + \lambda R(w^{(1)}) + \lambda R(w^{(2)}) \end{aligned}$$

- Computing gradients: $\frac{dL}{dw^{(1)}}$ and $\frac{dL}{dw^{(2)}}$
 - So that we can learn $w^{(1)}$ and $w^{(2)}$

Direct computation of gradients is extremely inefficient!!!

Computing Gradients

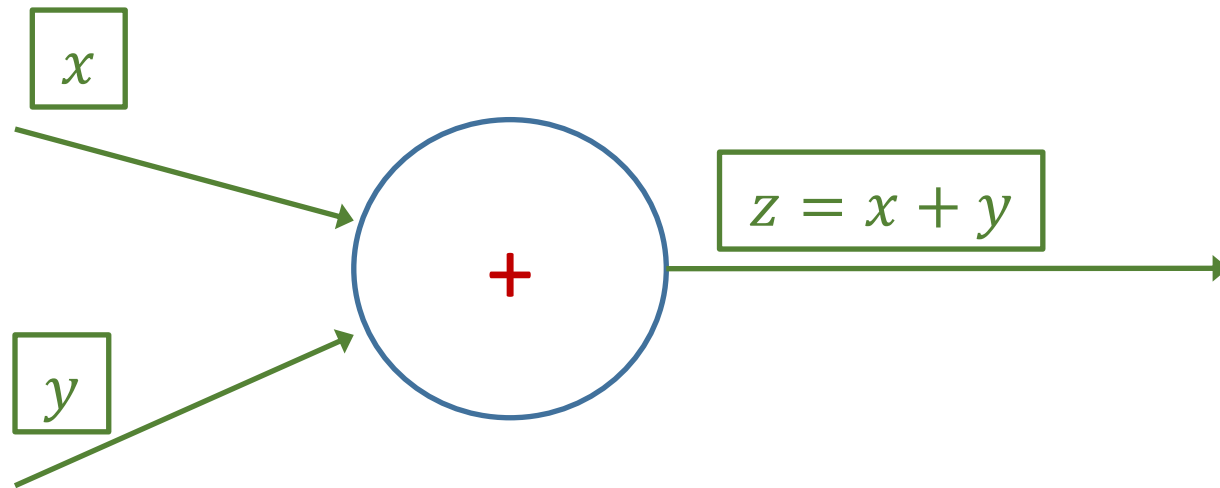
Computational Graphs
+
Backpropagation

Computational Graphs

- The descriptive language of deep learning models
- Functional description of the required computation
- Can be instantiated to do two types of computation:
 - Forward computation
 - Backward computation

Computational Graphs

- An acyclic directed graph
- Nodes: A node with an incoming edge is a function of that edge's tail node. This function can be basic arithmetic operations or any functions of which derivatives can be easily computed.
- Edges: An edge represents a function argument.



Computing Gradients: Computational Graphs + Backpropagation

- Key ideas: applying gradient chain rule to unroll gradients through hidden layers in a neural nets

- Chain rule: $\frac{df(g(x))}{dx} = \frac{df(g(x))}{dg(x)} \cdot \frac{dg(x)}{dx}$

- Example: $f(u) = u^2 + 2u$, $g(x) = 3x + 1$

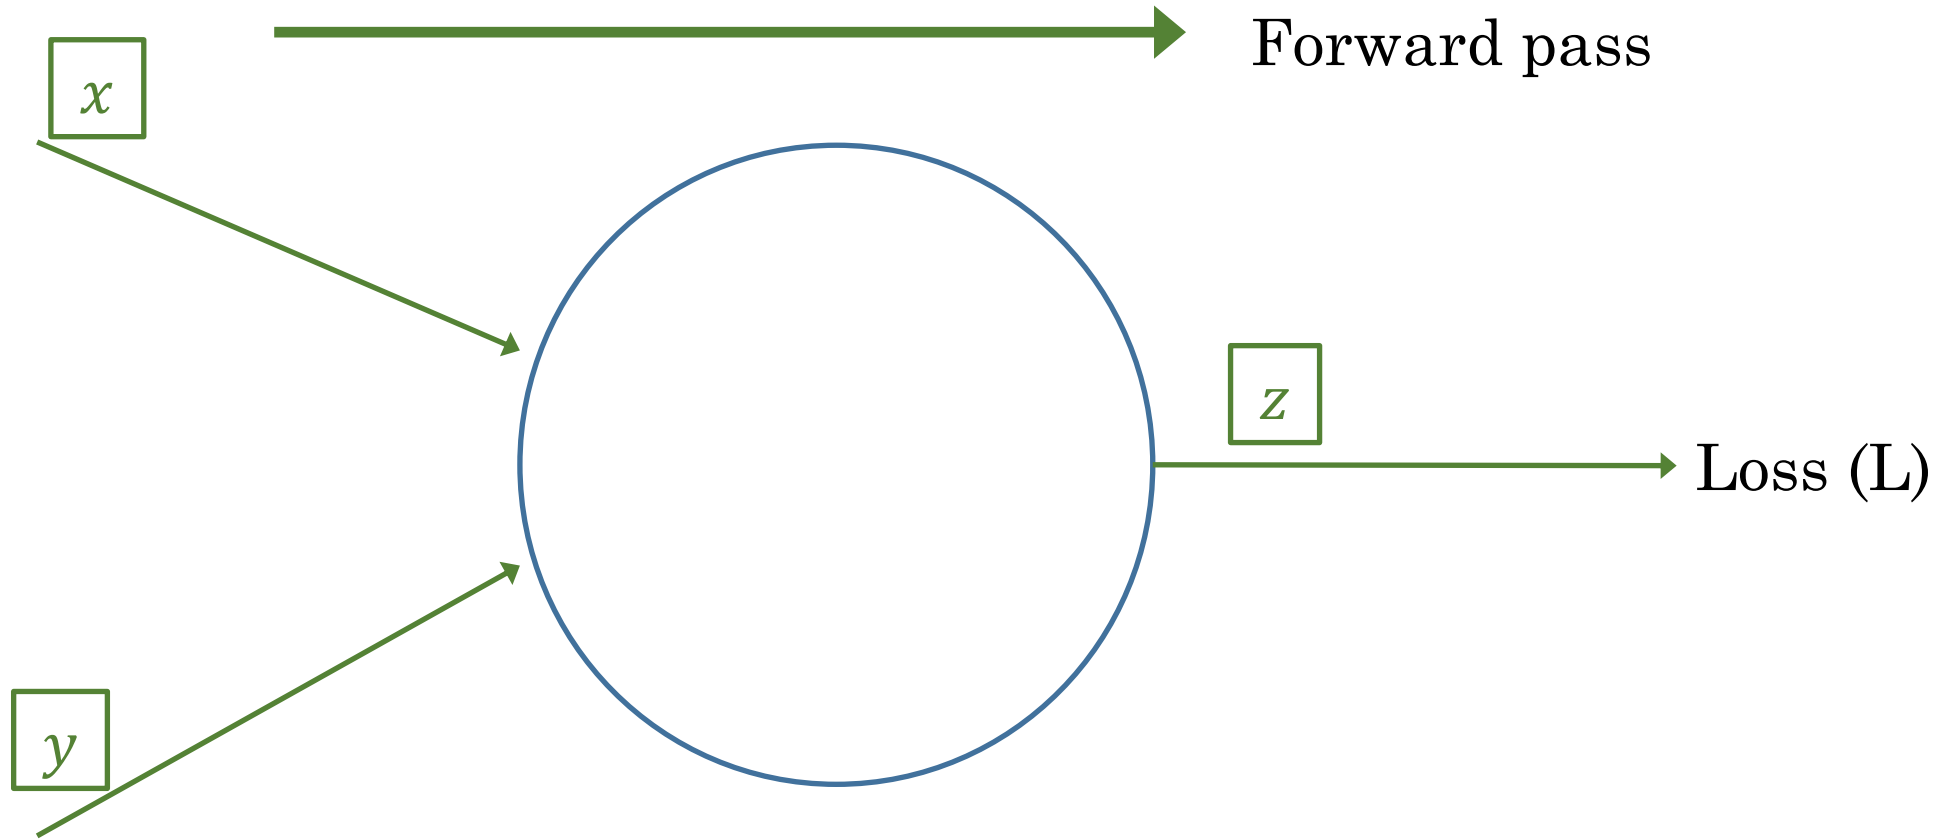
- Direct computation: $f(g(x)) = g(x)^2 + 2g(x) = (3x + 1)^2 + 2(3x + 1) = 9x^2 + 12x + 3 \rightarrow \frac{df(g(x))}{dx} = 18x + 12$ (Inefficient)

- Chain rule:

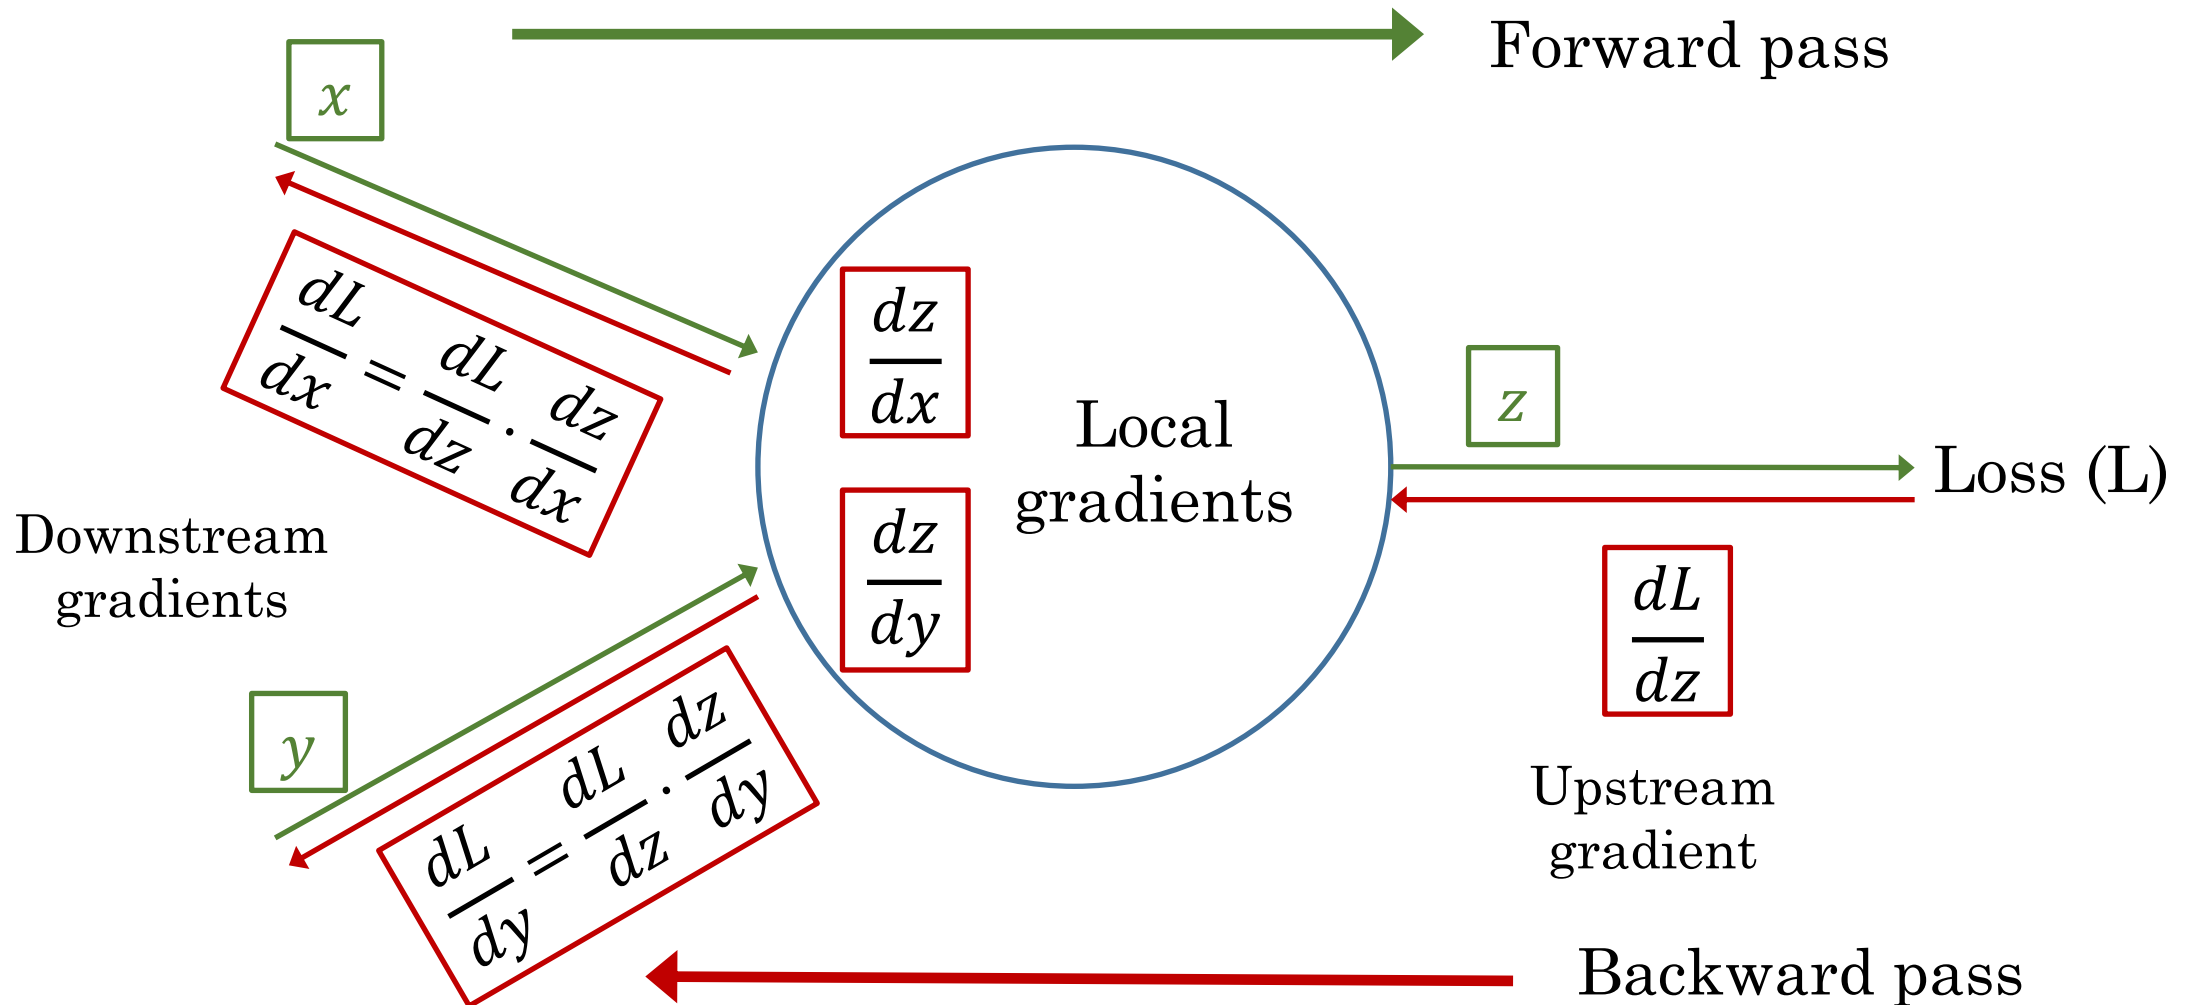
$$\frac{df(g(x))}{dx} = \frac{df(g(x))}{dg(x)} \cdot \frac{dg(x)}{dx} = (2g(x) + 2) \times 3 = (6x + 4) \times 3 = 18x + 12$$

(Much simpler)

Computational Graph + Backpropagation

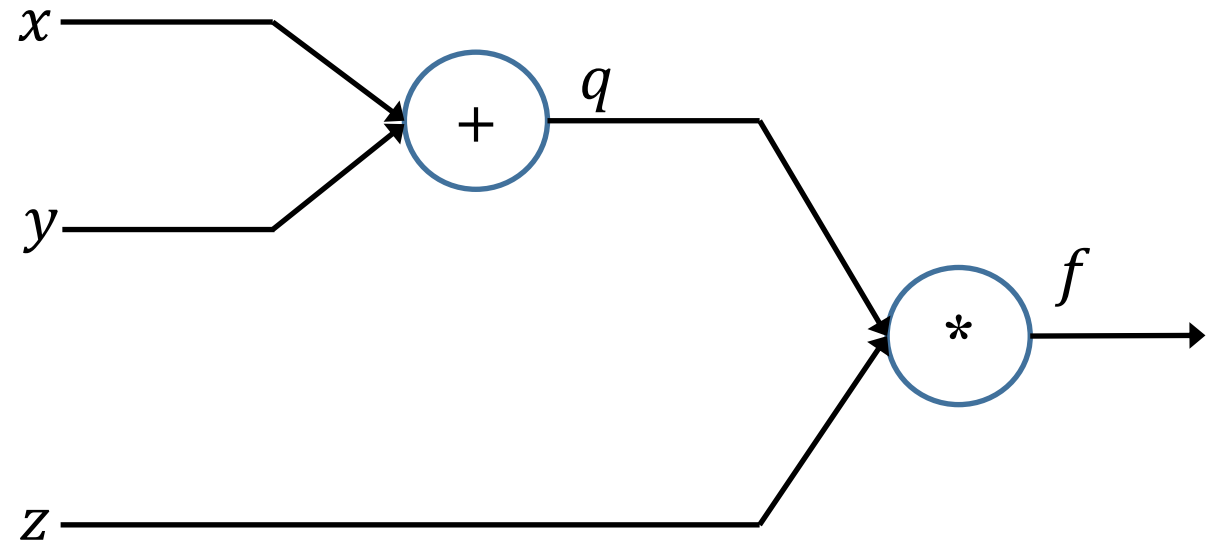


Computational Graph + Backpropagation



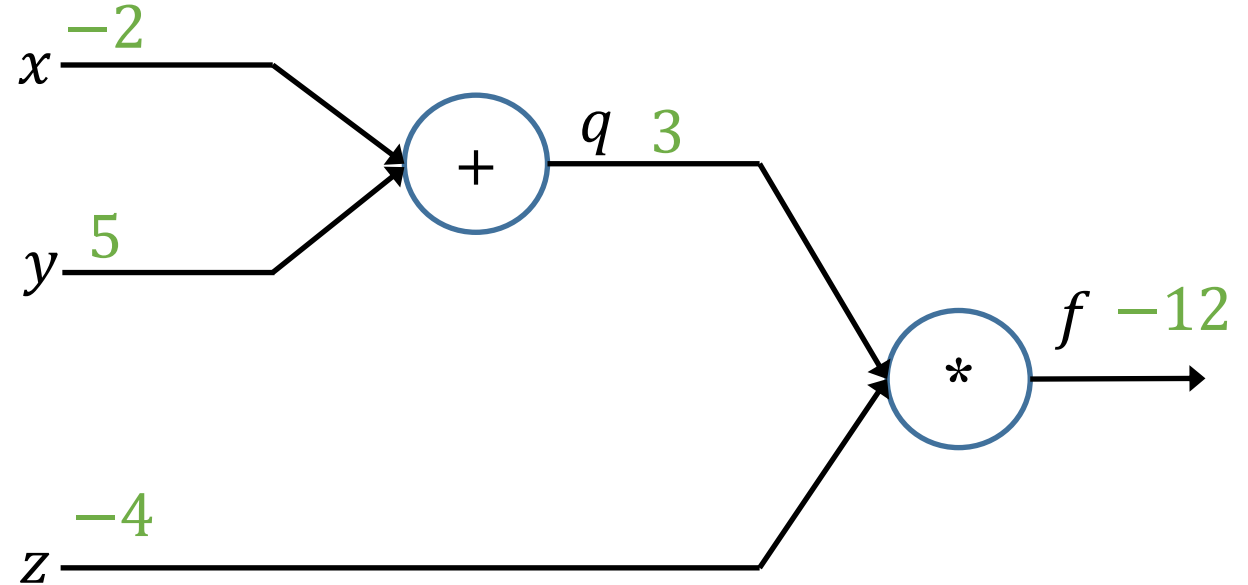
Example

- $q = x + y$
- $f = q * z$



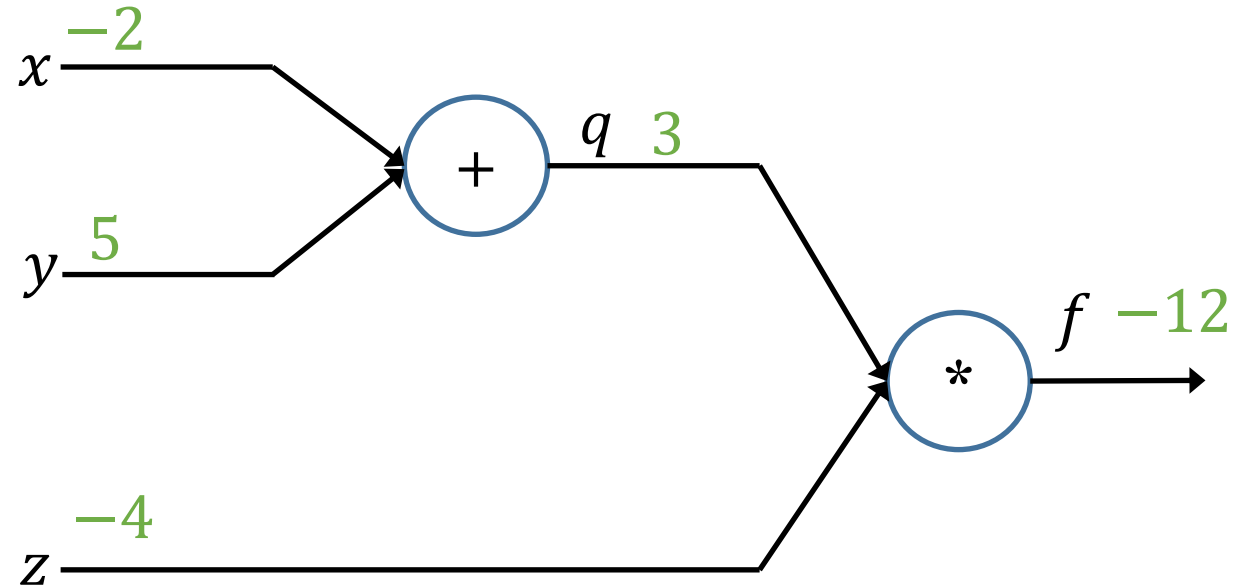
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



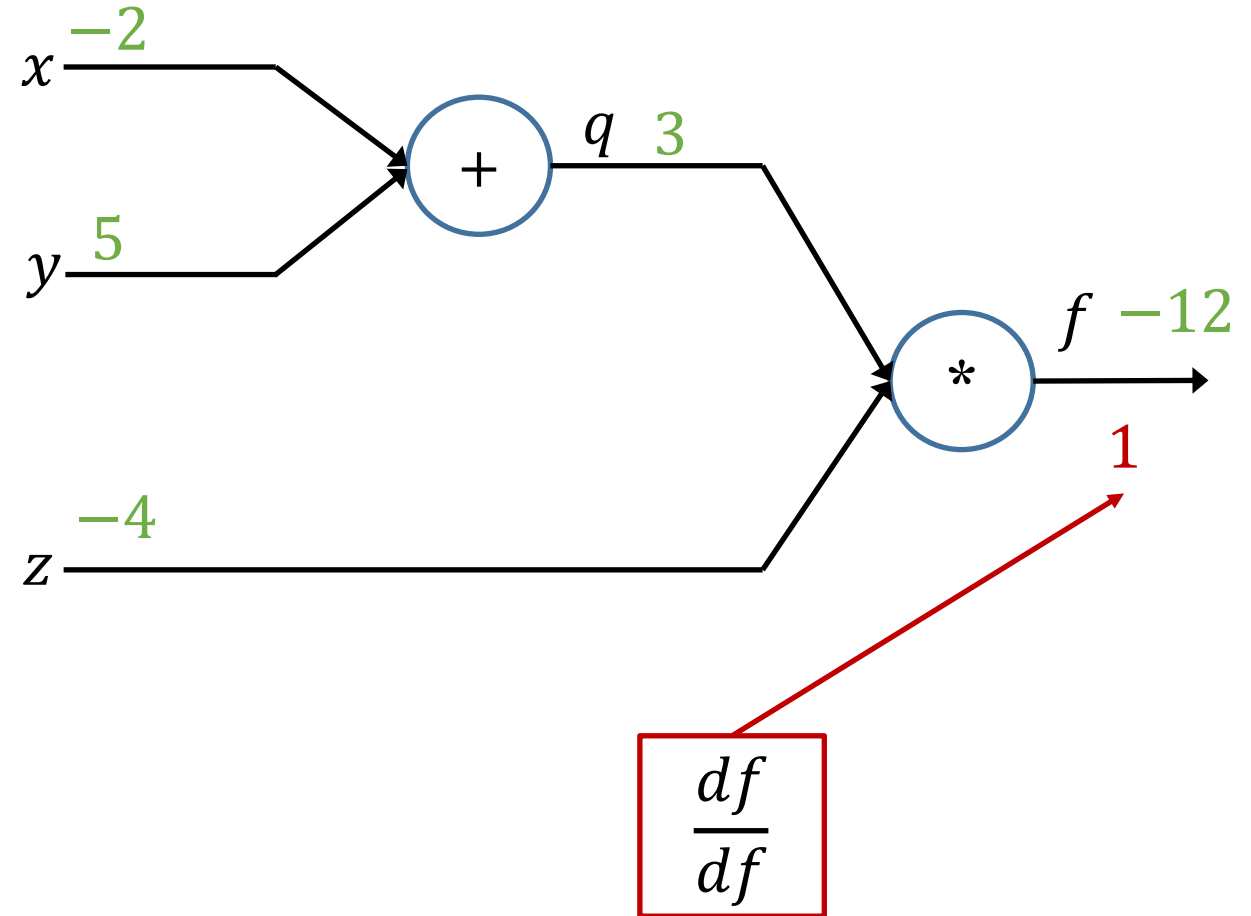
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



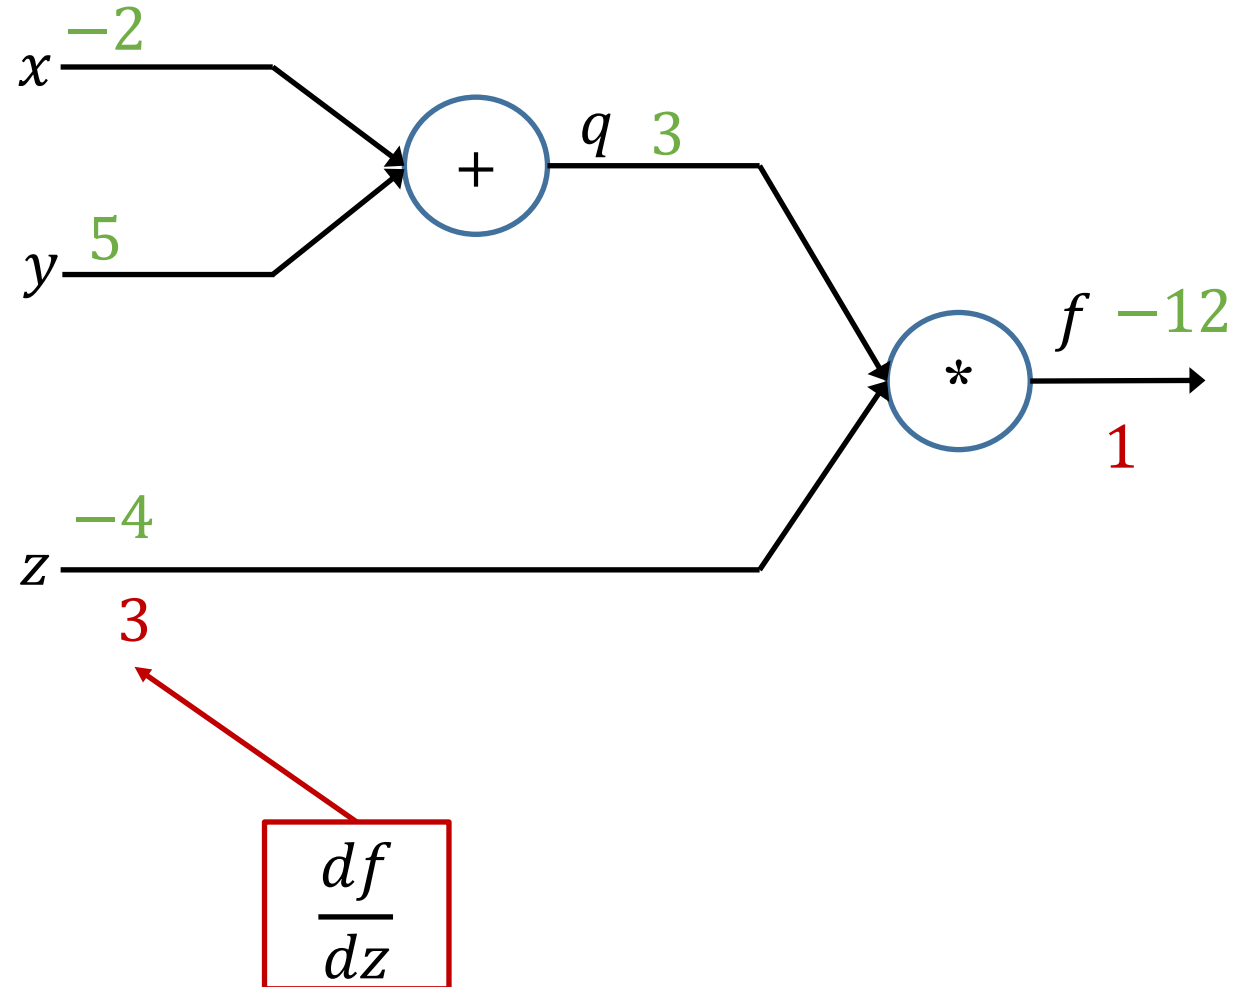
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



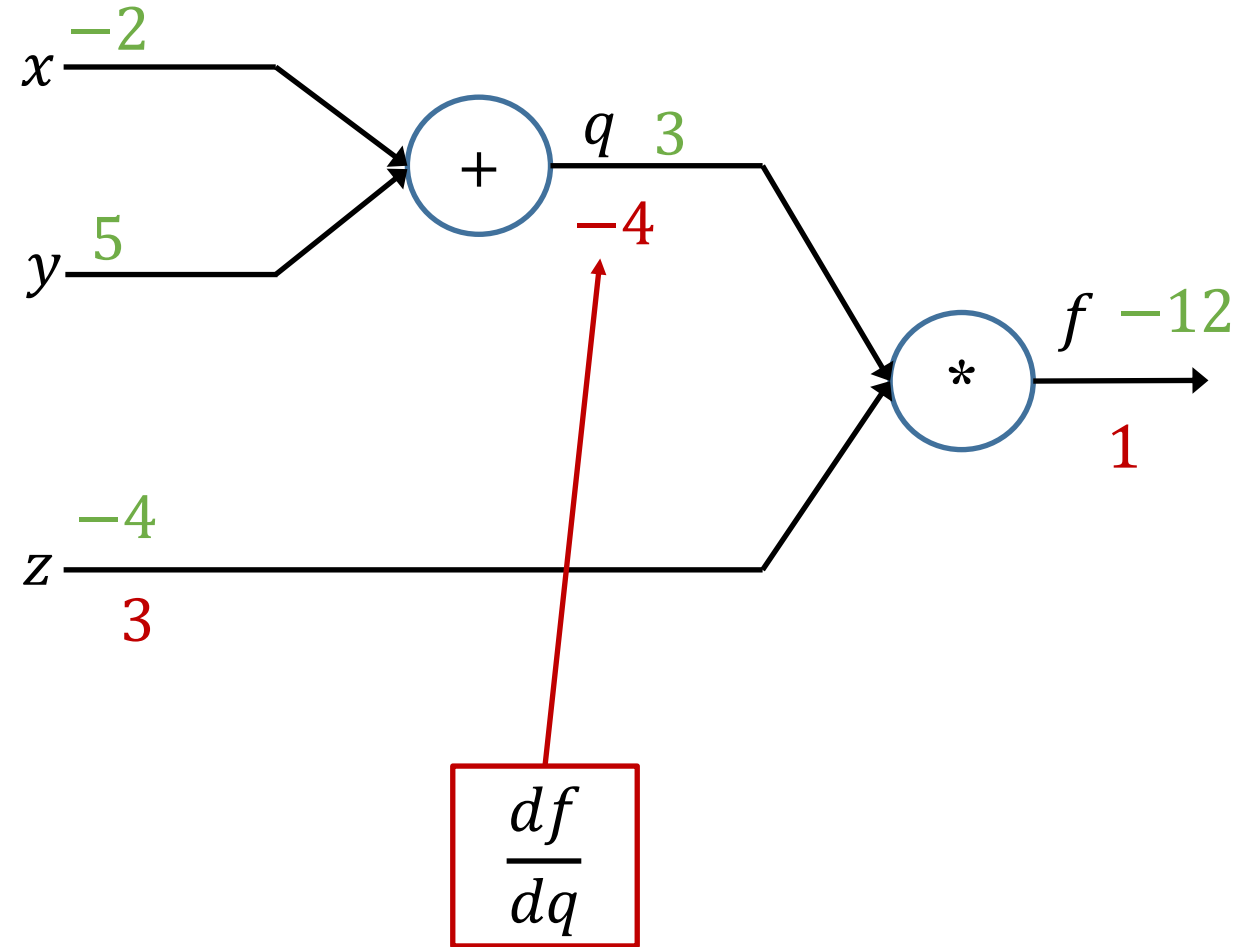
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



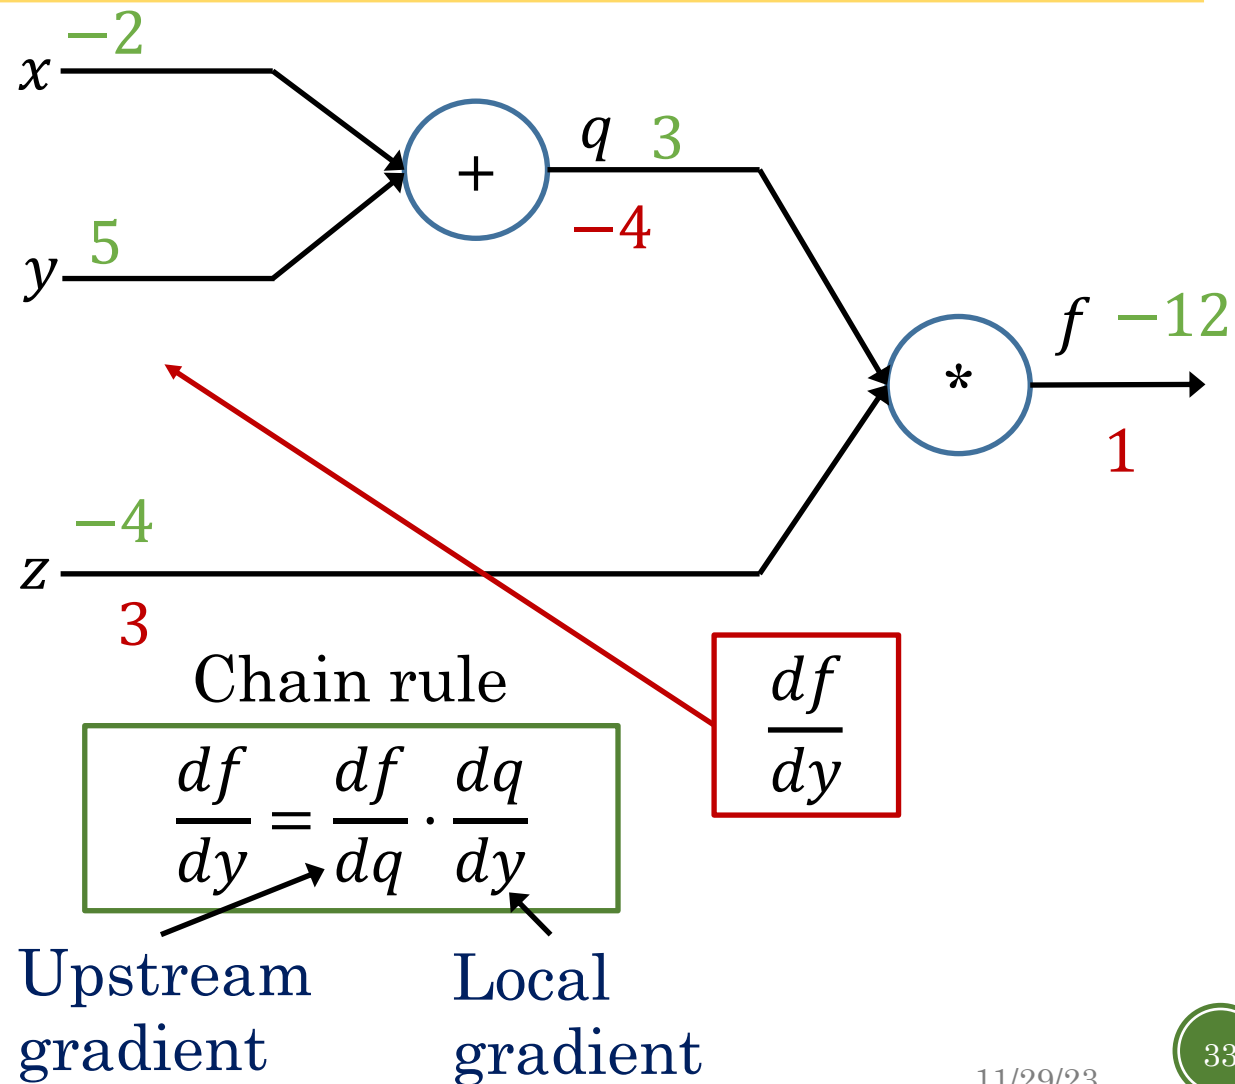
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



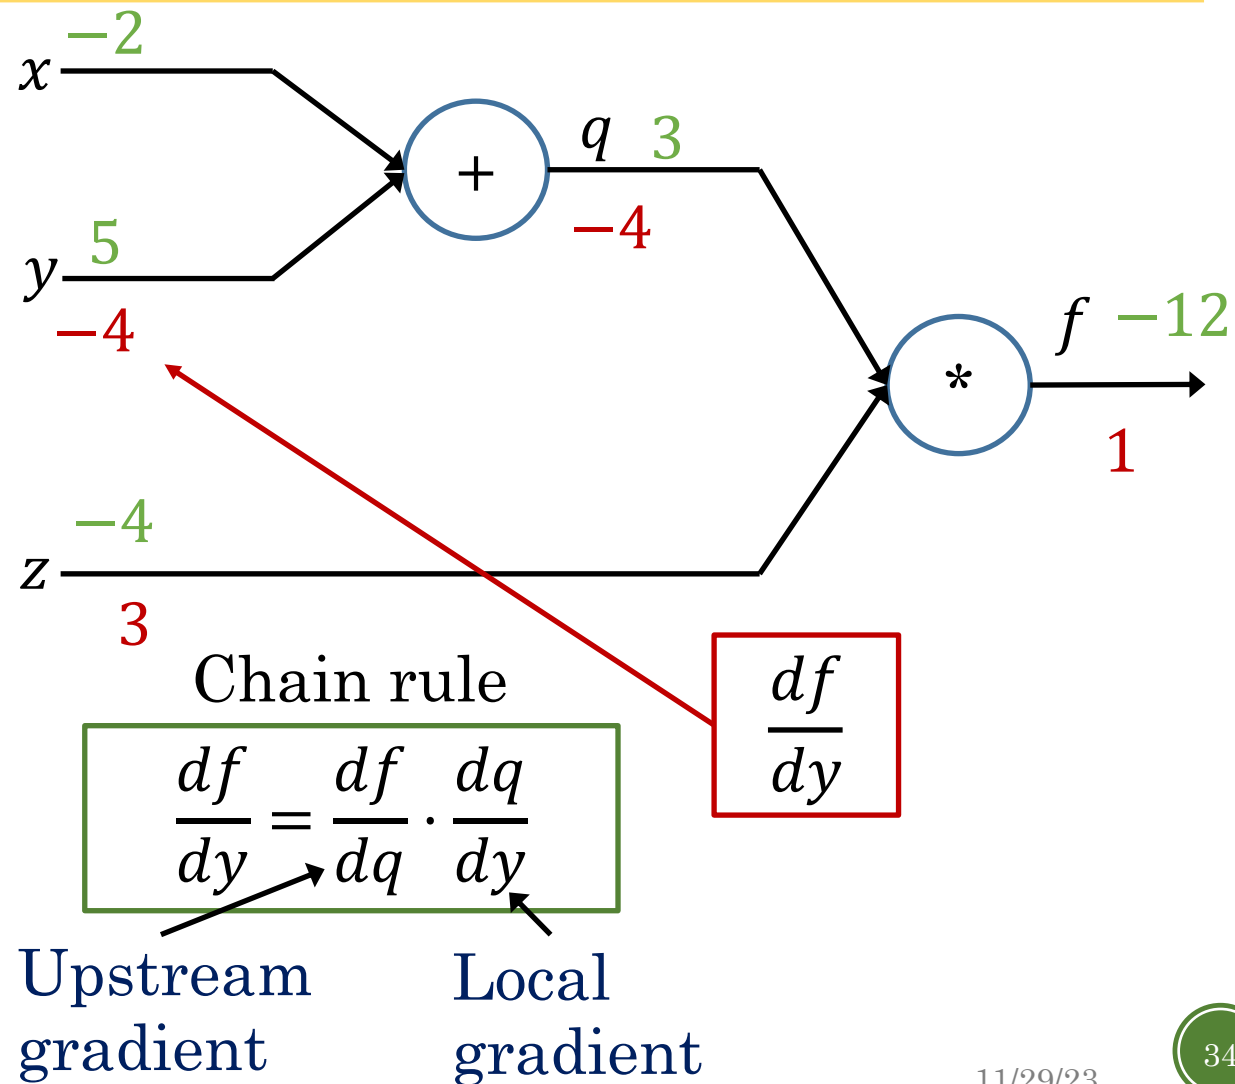
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



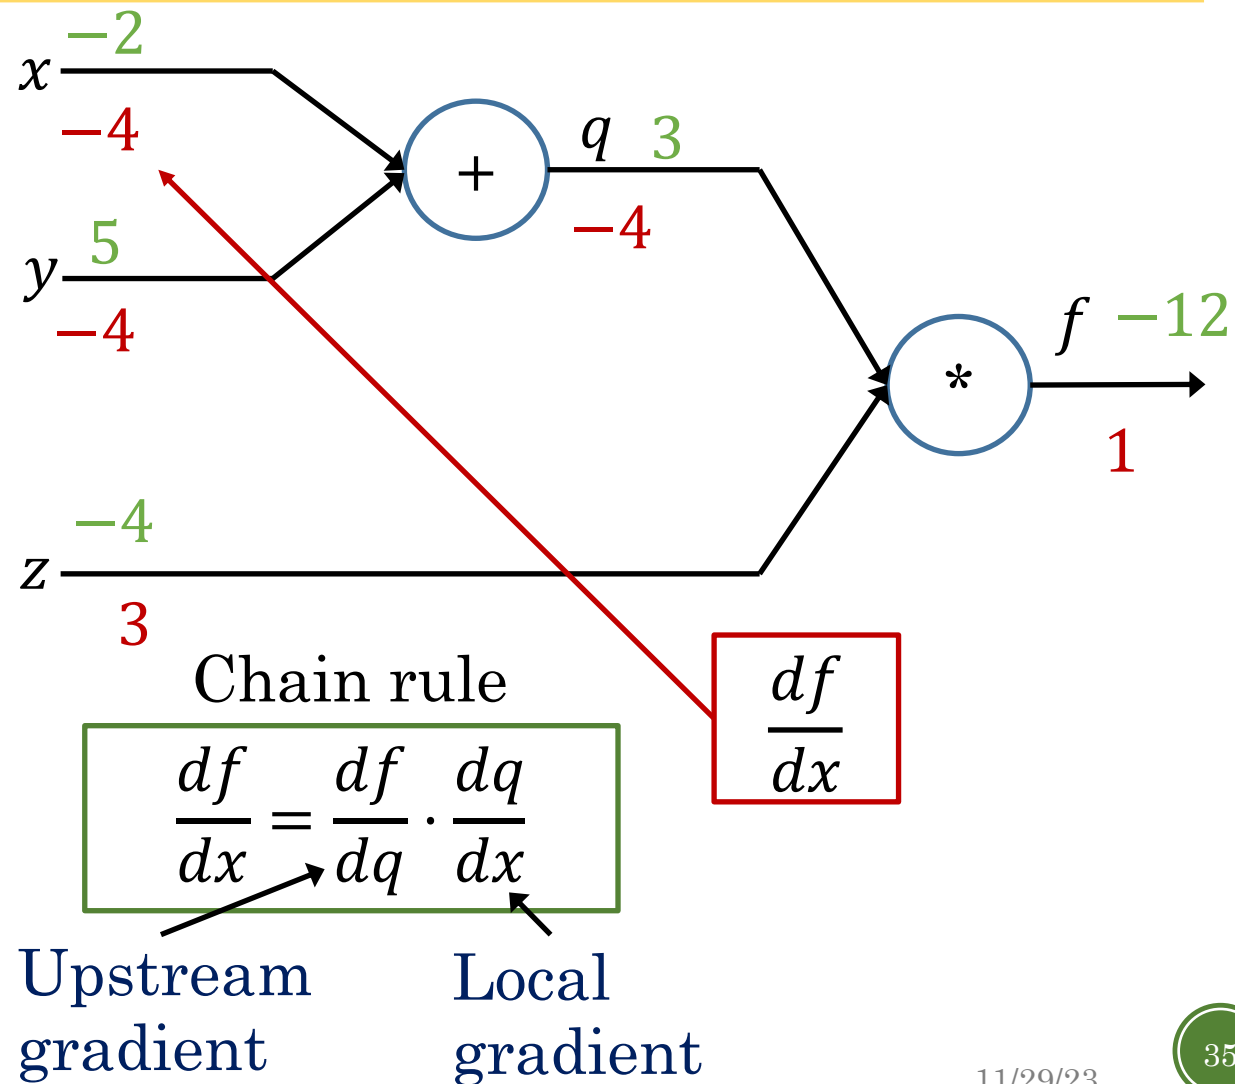
Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



Example

- $q = x + y$
- $f = q * z$
- Input: $x = -2, y = 5, z = -4$
- Goal: $\frac{df}{dx}, \frac{df}{dy}, \frac{df}{dz}$



Pytorch

- A Python-based scientific computing package of which goals are:
 - A replacement for NumPy to use the power of GPUs and other accelerators.
 - An automatic differentiation library that is useful to implement neural networks.
- GPU
 - A processor that has many smaller and more specialized cores
 - Has massive performance when a processing task can be divided up and processed across many cores.

Pytorch: Fundamental Concepts

- `torch.Tensor`: Tensors are the central data abstraction in PyTorch. Tensors are specialized data structure that are similar to arrays and matrices.
- `torch.autograd`: a built-in differentiation engine that supports automatic computation of gradient for any computational graph.
- `torch.nn.Module`: base class for all neural network modules.
- Pytorch tutorial:
 - https://pytorch.org/tutorials/beginner/pytorch_with_examples.html

Pytorch: Example

Create data



```
1 import torch
2 import math
3
4 dtype = torch.float
5 device = torch.device("cpu")
6
7 x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8 y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36 print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```

Pytorch: Example

Create random weights



```
1  import torch
2  import math
3
4  dtype = torch.float
5  device = torch.device("cpu")
6
7  x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8  y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36     print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```

Pytorch: Example

Forward pass



```
1 import torch
2 import math
3
4 dtype = torch.float
5 device = torch.device("cpu")
6
7 x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8 y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36     print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```


Pytorch: Example

Backward pass



```
1 import torch
2 import math
3
4 dtype = torch.float
5 device = torch.device("cpu")
6
7 x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8 y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36     print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```

Pytorch: Example


Gradient descent



```
1 import torch
2 import math
3
4 dtype = torch.float
5 device = torch.device("cpu")
6
7 x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8 y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36     print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```

Pytorch: Example

Reset gradient before
next round of forward-
backward pass



```
1 import torch
2 import math
3
4 dtype = torch.float
5 device = torch.device("cpu")
6
7 x = torch.linspace(-math.pi, math.pi, 2000, device = device, dtype = dtype)
8 y = torch.sin(x)
9
10 a = torch.randn((), device = device, dtype = dtype, requires_grad=True)
11 b = torch.randn((), device = device, dtype = dtype, requires_grad=True)
12 c = torch.randn((), device = device, dtype = dtype, requires_grad=True)
13 d = torch.randn((), device = device, dtype = dtype, requires_grad=True)
14
15 learning_rate = 1e-6
16
17 for t in range(2000):
18     y_pred = a + b * x + c * x ** 2 + d * x ** 3
19     loss = (y_pred - y).pow(2).sum()
20
21     if t % 100 == 99:
22         print(t, loss.item())
23
24     loss.backward()
25     with torch.no_grad():
26         a -= learning_rate * a.grad
27         b -= learning_rate * b.grad
28         c -= learning_rate * c.grad
29         d -= learning_rate * d.grad
30
31     a.grad = None
32     b.grad = None
33     c.grad = None
34     d.grad = None
35
36     print(f'Result: y = {a.item()} + {b.item()} x + {c.item()} x^2 + {d.item()} x^3')
```

Pytorch: Neural Network

Create data



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24     model.zero_grad()
25
26     loss.backward()
27
28     with torch.no_grad():
29         for param in model.parameters():
30             param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```

Pytorch: Neural Network

Create predictive
model



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24     model.zero_grad()
25
26     loss.backward()
27
28     with torch.no_grad():
29         for param in model.parameters():
30             param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```

Pytorch: Neural Network

Define loss
function



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24     model.zero_grad()
25
26     loss.backward()
27
28     with torch.no_grad():
29         for param in model.parameters():
30             param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```

Pytorch: Neural Network

Forward pass



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24         model.zero_grad()
25
26         loss.backward()
27
28         with torch.no_grad():
29             for param in model.parameters():
30                 param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```

Pytorch: Neural Network

Reset gradient and
run backward pass



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24     model.zero_grad()
25
26     loss.backward()
27
28     with torch.no_grad():
29         for param in model.parameters():
30             param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```


Pytorch: Neural Network

Run gradient
descent



```
1 import torch
2 import math
3
4 x = torch.linspace(-math.pi, math.pi, 2000)
5 y = torch.sin(x)
6
7 p = torch.tensor([1, 2, 3])
8 xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-6
18 for t in range(2000):
19     y_pred = model(xx)
20
21     loss = loss_fn(y_pred, y)
22     if t % 100 == 99:
23         print(t, loss.item())
24     model.zero_grad()
25
26     loss.backward()
27
28     with torch.no_grad():
29         for param in model.parameters():
30             param -= learning_rate * param.grad
31
32 linear_layer = model[0]
```

Pytorch: Optimizer

Choose an optimizer



```
1  import torch
2  import math
3
4  x = torch.linspace(-math.pi, math.pi, 2000)
5  y = torch.sin(x)
6
7  p = torch.tensor([1, 2, 3])
8  xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-3
18 optimizer = torch.optim.RMSprop(model.parameters(), lr=learning_rate)
19 for t in range(2000):
20
21     y_pred = model(xx)
22     loss = loss_fn(y_pred, y)
23
24     if t % 100 == 99:
25         print(t, loss.item())
26     model.zero_grad()
27
28     loss.backward()
29
30     optimizer.step()
```

Pytorch: Optimizer

Update weights


Thanh H. Nguyen



```
1  import torch
2  import math
3
4  x = torch.linspace(-math.pi, math.pi, 2000)
5  y = torch.sin(x)
6
7  p = torch.tensor([1, 2, 3])
8  xx = x.unsqueeze(-1).pow(p)
9
10 model = torch.nn.Sequential(
11     torch.nn.Linear(3, 1),
12     torch.nn.Flatten(0, 1)
13 )
14
15 loss_fn = torch.nn.MSELoss(reduction='sum')
16
17 learning_rate = 1e-3
18 optimizer = torch.optim.RMSprop(model.parameters(), lr=learning_rate)
19 for t in range(2000):
20     y_pred = model(xx)
21     loss = loss_fn(y_pred, y)
22
23     if t % 100 == 99:
24         print(t, loss.item())
25     model.zero_grad()
26
27     loss.backward()
28
29     optimizer.step()
```

Pytorch: Define a New Module


Define a neural network
model as a module



```
1 import torch
2
3
4 class Net(torch.nn.Module):
5     def __init__(self, d_in, d_hidden, d_out):
6         super(Net, self).__init__()
7         self.ln1 = torch.nn.Linear(d_in, d_hidden)
8         self.ln2 = torch.nn.Linear(d_hidden, d_out)
9
10    def forward(self, x):
11        h_relu = self.ln1(x).clamp(min=0)
12        y_pred = self.ln2(h_relu)
13        return y_pred
14
15
16 train_size, d_in, d_hidden, d_out = 64, 1000, 100, 10
17 x = torch.randn(train_size, d_in)
18 y = torch.randn(train_size, d_out)
19
20 model = Net(d_in, d_hidden, d_out)
21 optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
22 loss_fn = torch.nn.MSELoss(reduction='sum')
23
24 for t in range(500):
25     y_pred = model(x)
26     loss = loss_fn(y_pred, y)
27
28     if t % 50 == 49:
29         print(t, loss.item())
30
31     model.zero_grad()
32     loss.backward()
33
34     optimizer.step()
35
```

Pytorch: Define a New Module

Construct and train a
model



```
1 import torch
2
3
4 class Net(torch.nn.Module):
5     def __init__(self, d_in, d_hidden, d_out):
6         super(Net, self).__init__()
7         self.ln1 = torch.nn.Linear(d_in, d_hidden)
8         self.ln2 = torch.nn.Linear(d_hidden, d_out)
9
10    def forward(self, x):
11        h_relu = self.ln1(x).clamp(min=0)
12        y_pred = self.ln2(h_relu)
13        return y_pred
14
15
16 train_size, d_in, d_hidden, d_out = 64, 1000, 100, 10
17 x = torch.randn(train_size, d_in)
18 y = torch.randn(train_size, d_out)
19
20 model = Net(d_in, d_hidden, d_out)
21 optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
22 loss_fn = torch.nn.MSELoss(reduction='sum')
23
24 for t in range(500):
25     y_pred = model(x)
26     loss = loss_fn(y_pred, y)
27
28     if t % 50 == 49:
29         print(t, loss.item())
30
31     model.zero_grad()
32     loss.backward()
33
34     optimizer.step()
35
```